

Return-Oriented Programming

Nathan Chong
22 March 2016

This talk

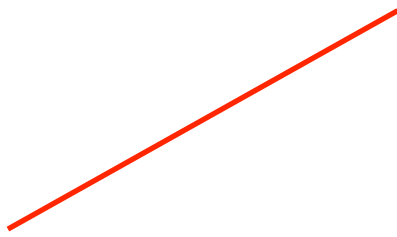
- Study a stack-smash exploit
- Defences (particularly XN)
- Attacks
 - Return-to-Libc
 - Return-Oriented Programming

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```

Source: <http://blog.techorganic.com/2015/04/10/64-bit-linux-stack-smashing-tutorial-part-1/>

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```

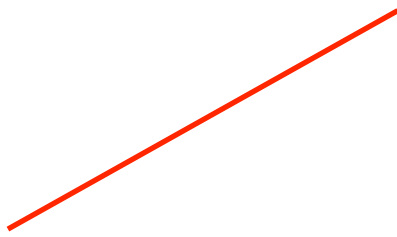
Read 400 bytes from
STDIN (fd 0) into buf



Bug found: Job done! End-of-Talk :-)

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```

Read 400 bytes from
STDIN (fd 0) into buf



Buffer-overflow bug allows stack-smash exploit

i.e., subvert the program's control flow to
a payload of the attacker's choosing

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```



```
gcc -masm=intel -S \  
    -fno-stack-protector
```



revisit this later

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```



```
gcc -masm=intel -S \  
    -fno-stack-protector
```

```
vuln:  
    push rbp  
    mov rbp, rsp  
    sub rsp, 96  
    lea rax, [rbp-96]  
    mov edx, 400  
    mov rsi, rax  
    mov edi, 0  
    call read  
  
    // ...  
  
    leave  
    ret
```

```
int vuln() {  
    char buf[80];  
    int r;  
    r = read(0, buf, 400);  
    // ...  
    return 0;  
}
```

function entry

argument handling
and read() call

function return

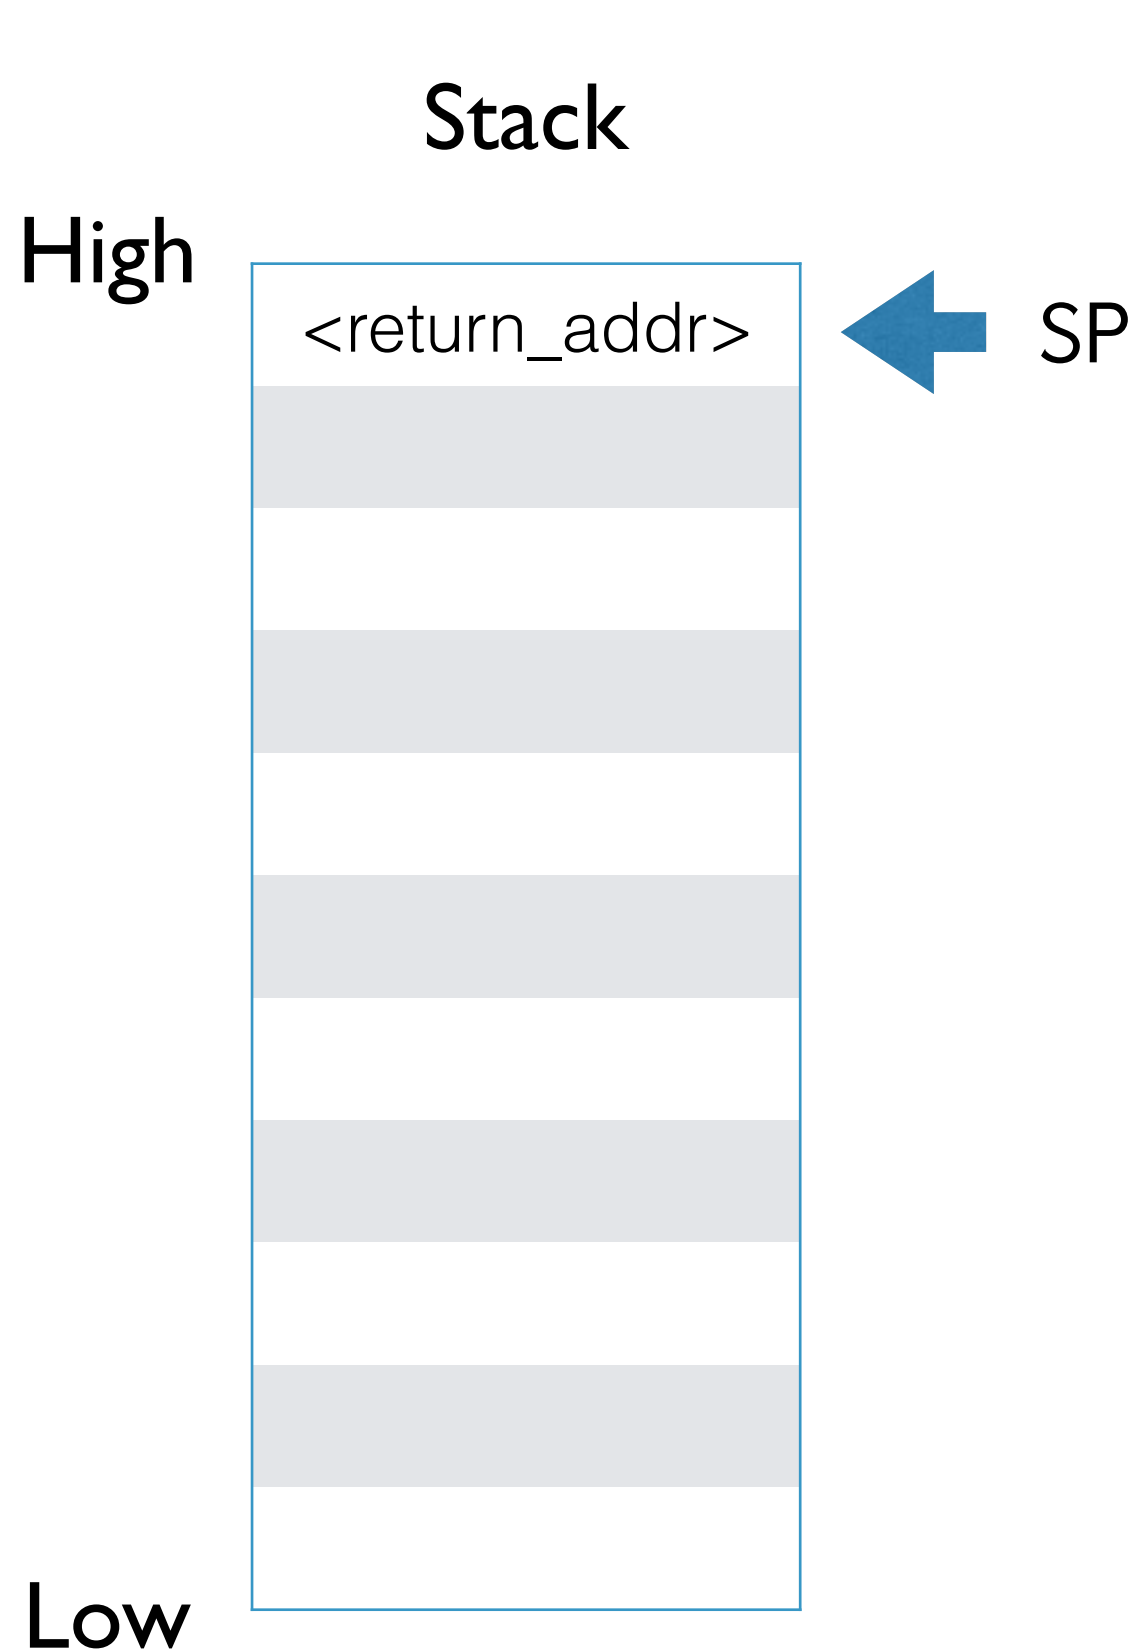
vuln:

```
push rbp  
mov rbp, rsp  
sub rsp, 96
```

```
lea rax, [rbp-96]  
mov edx, 400  
mov rsi, rax  
mov edi, 0  
call read
```

// ...

```
leave  
ret
```

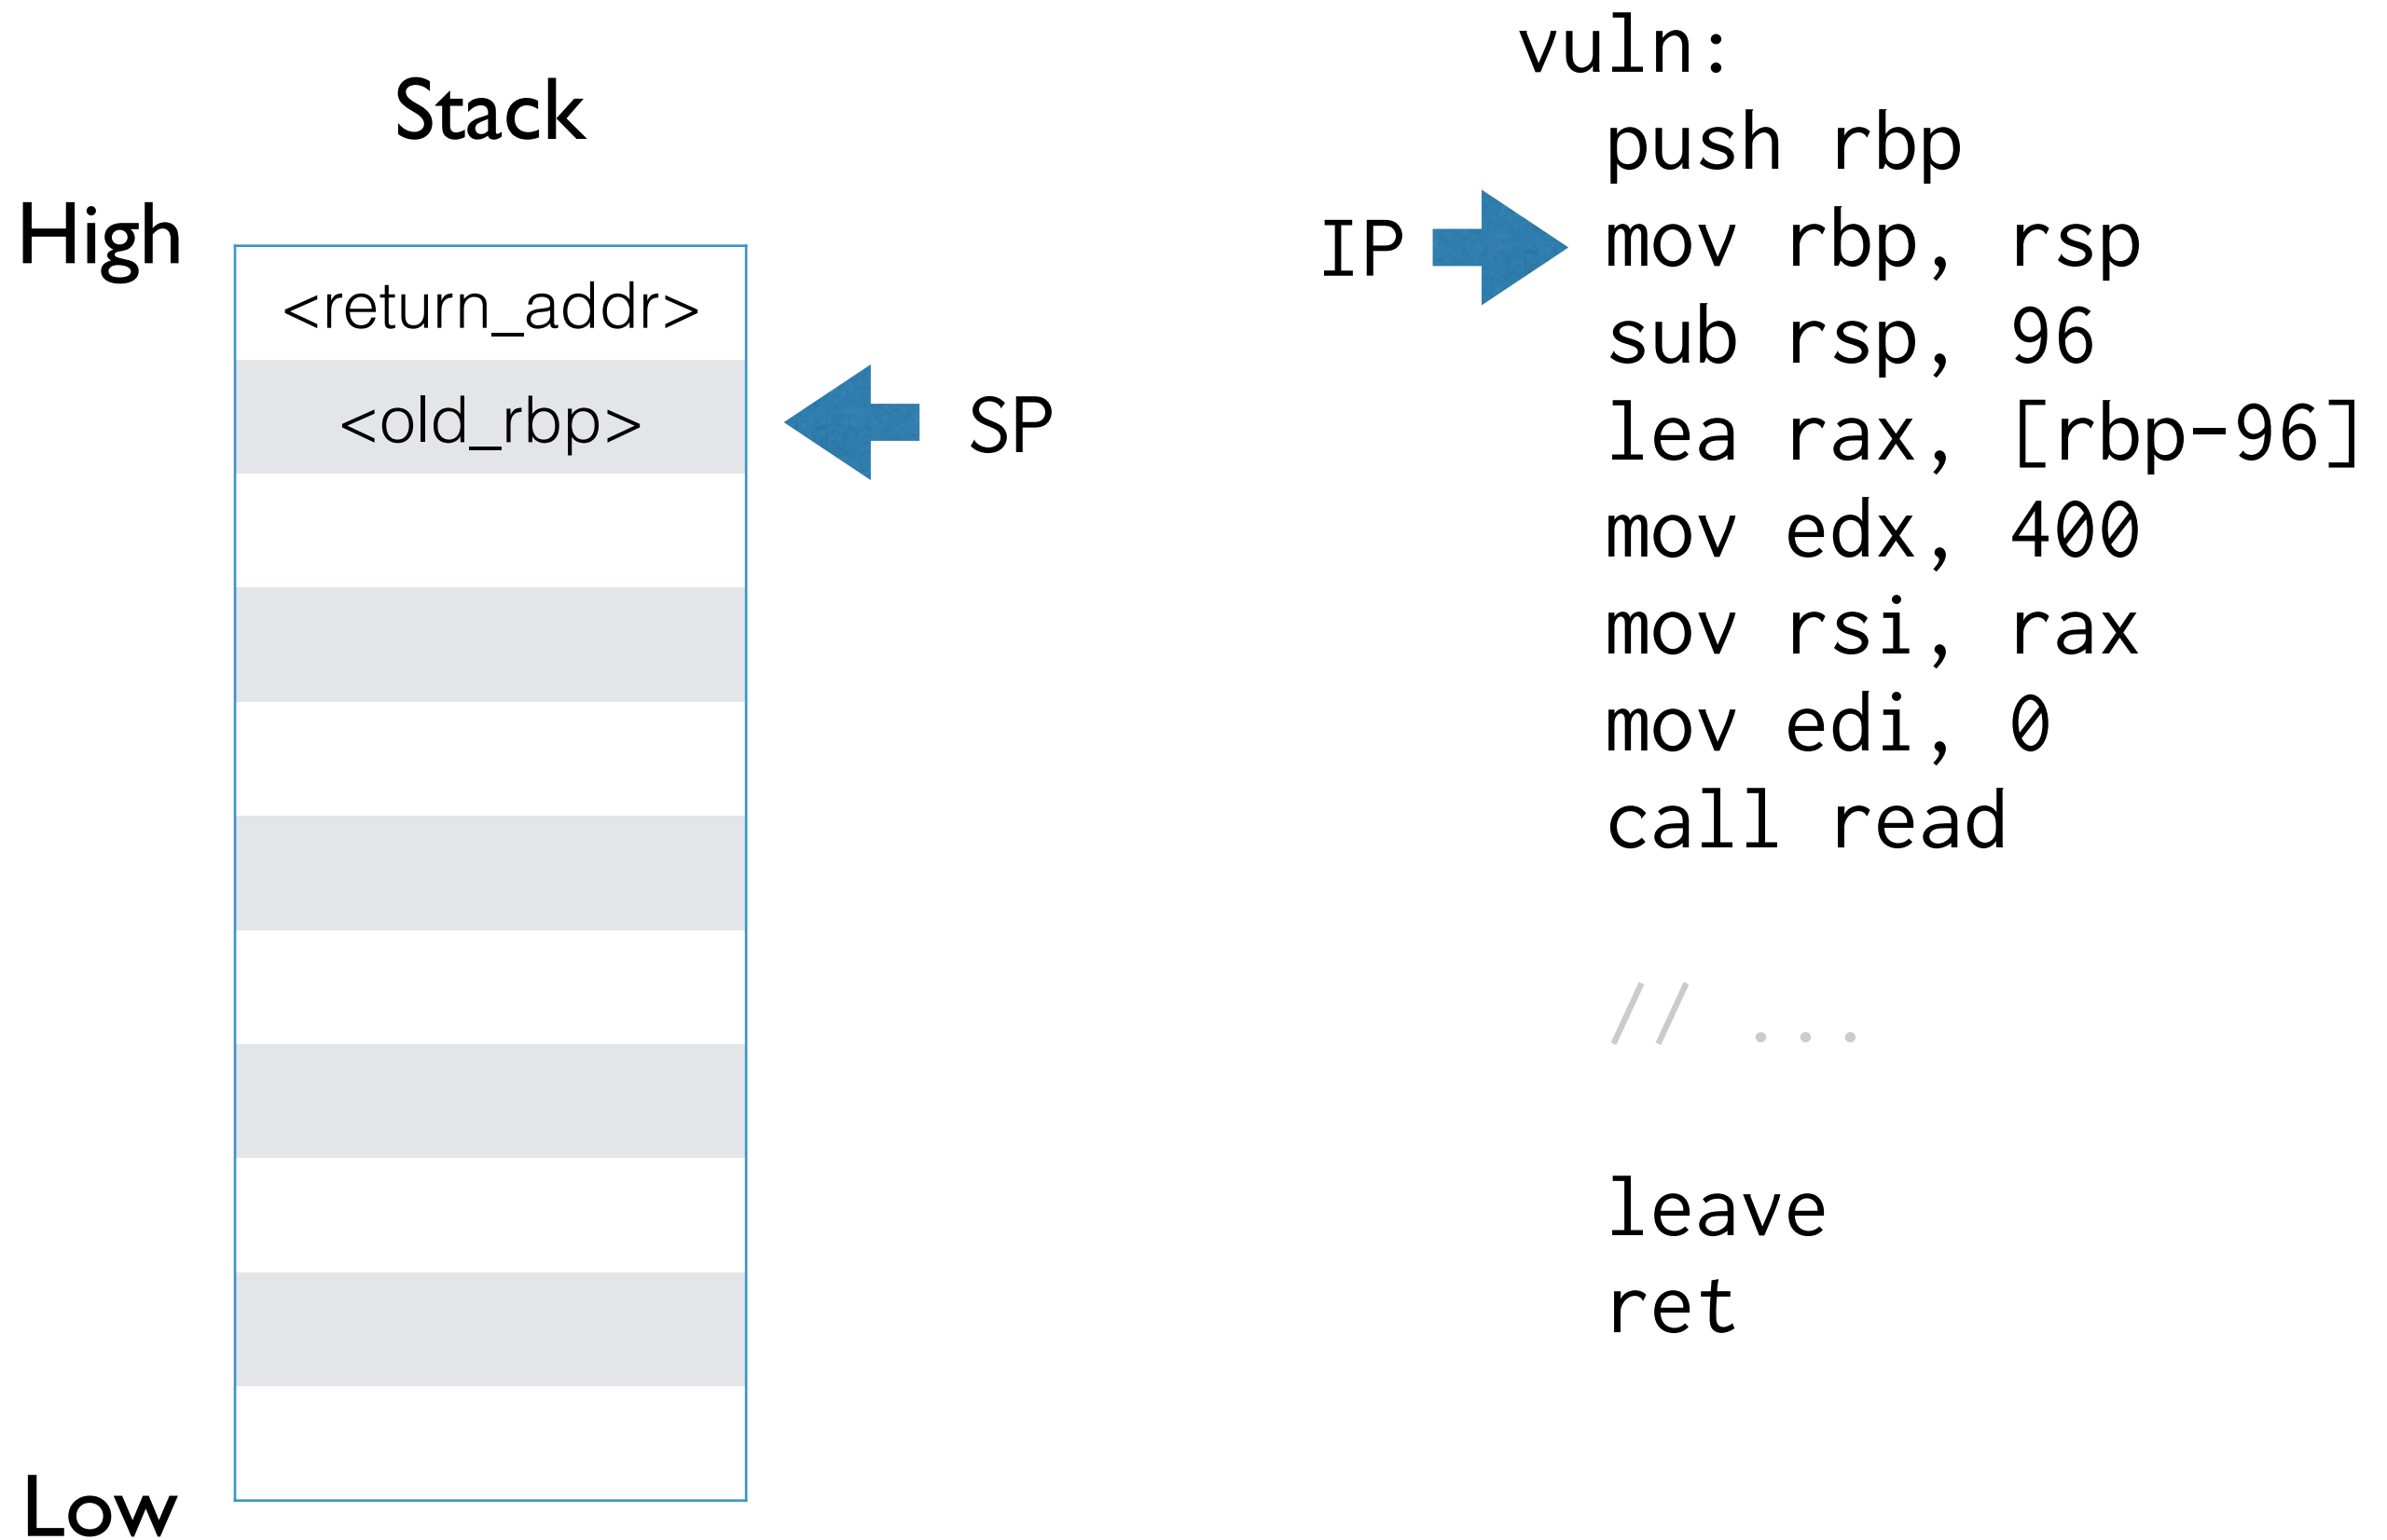



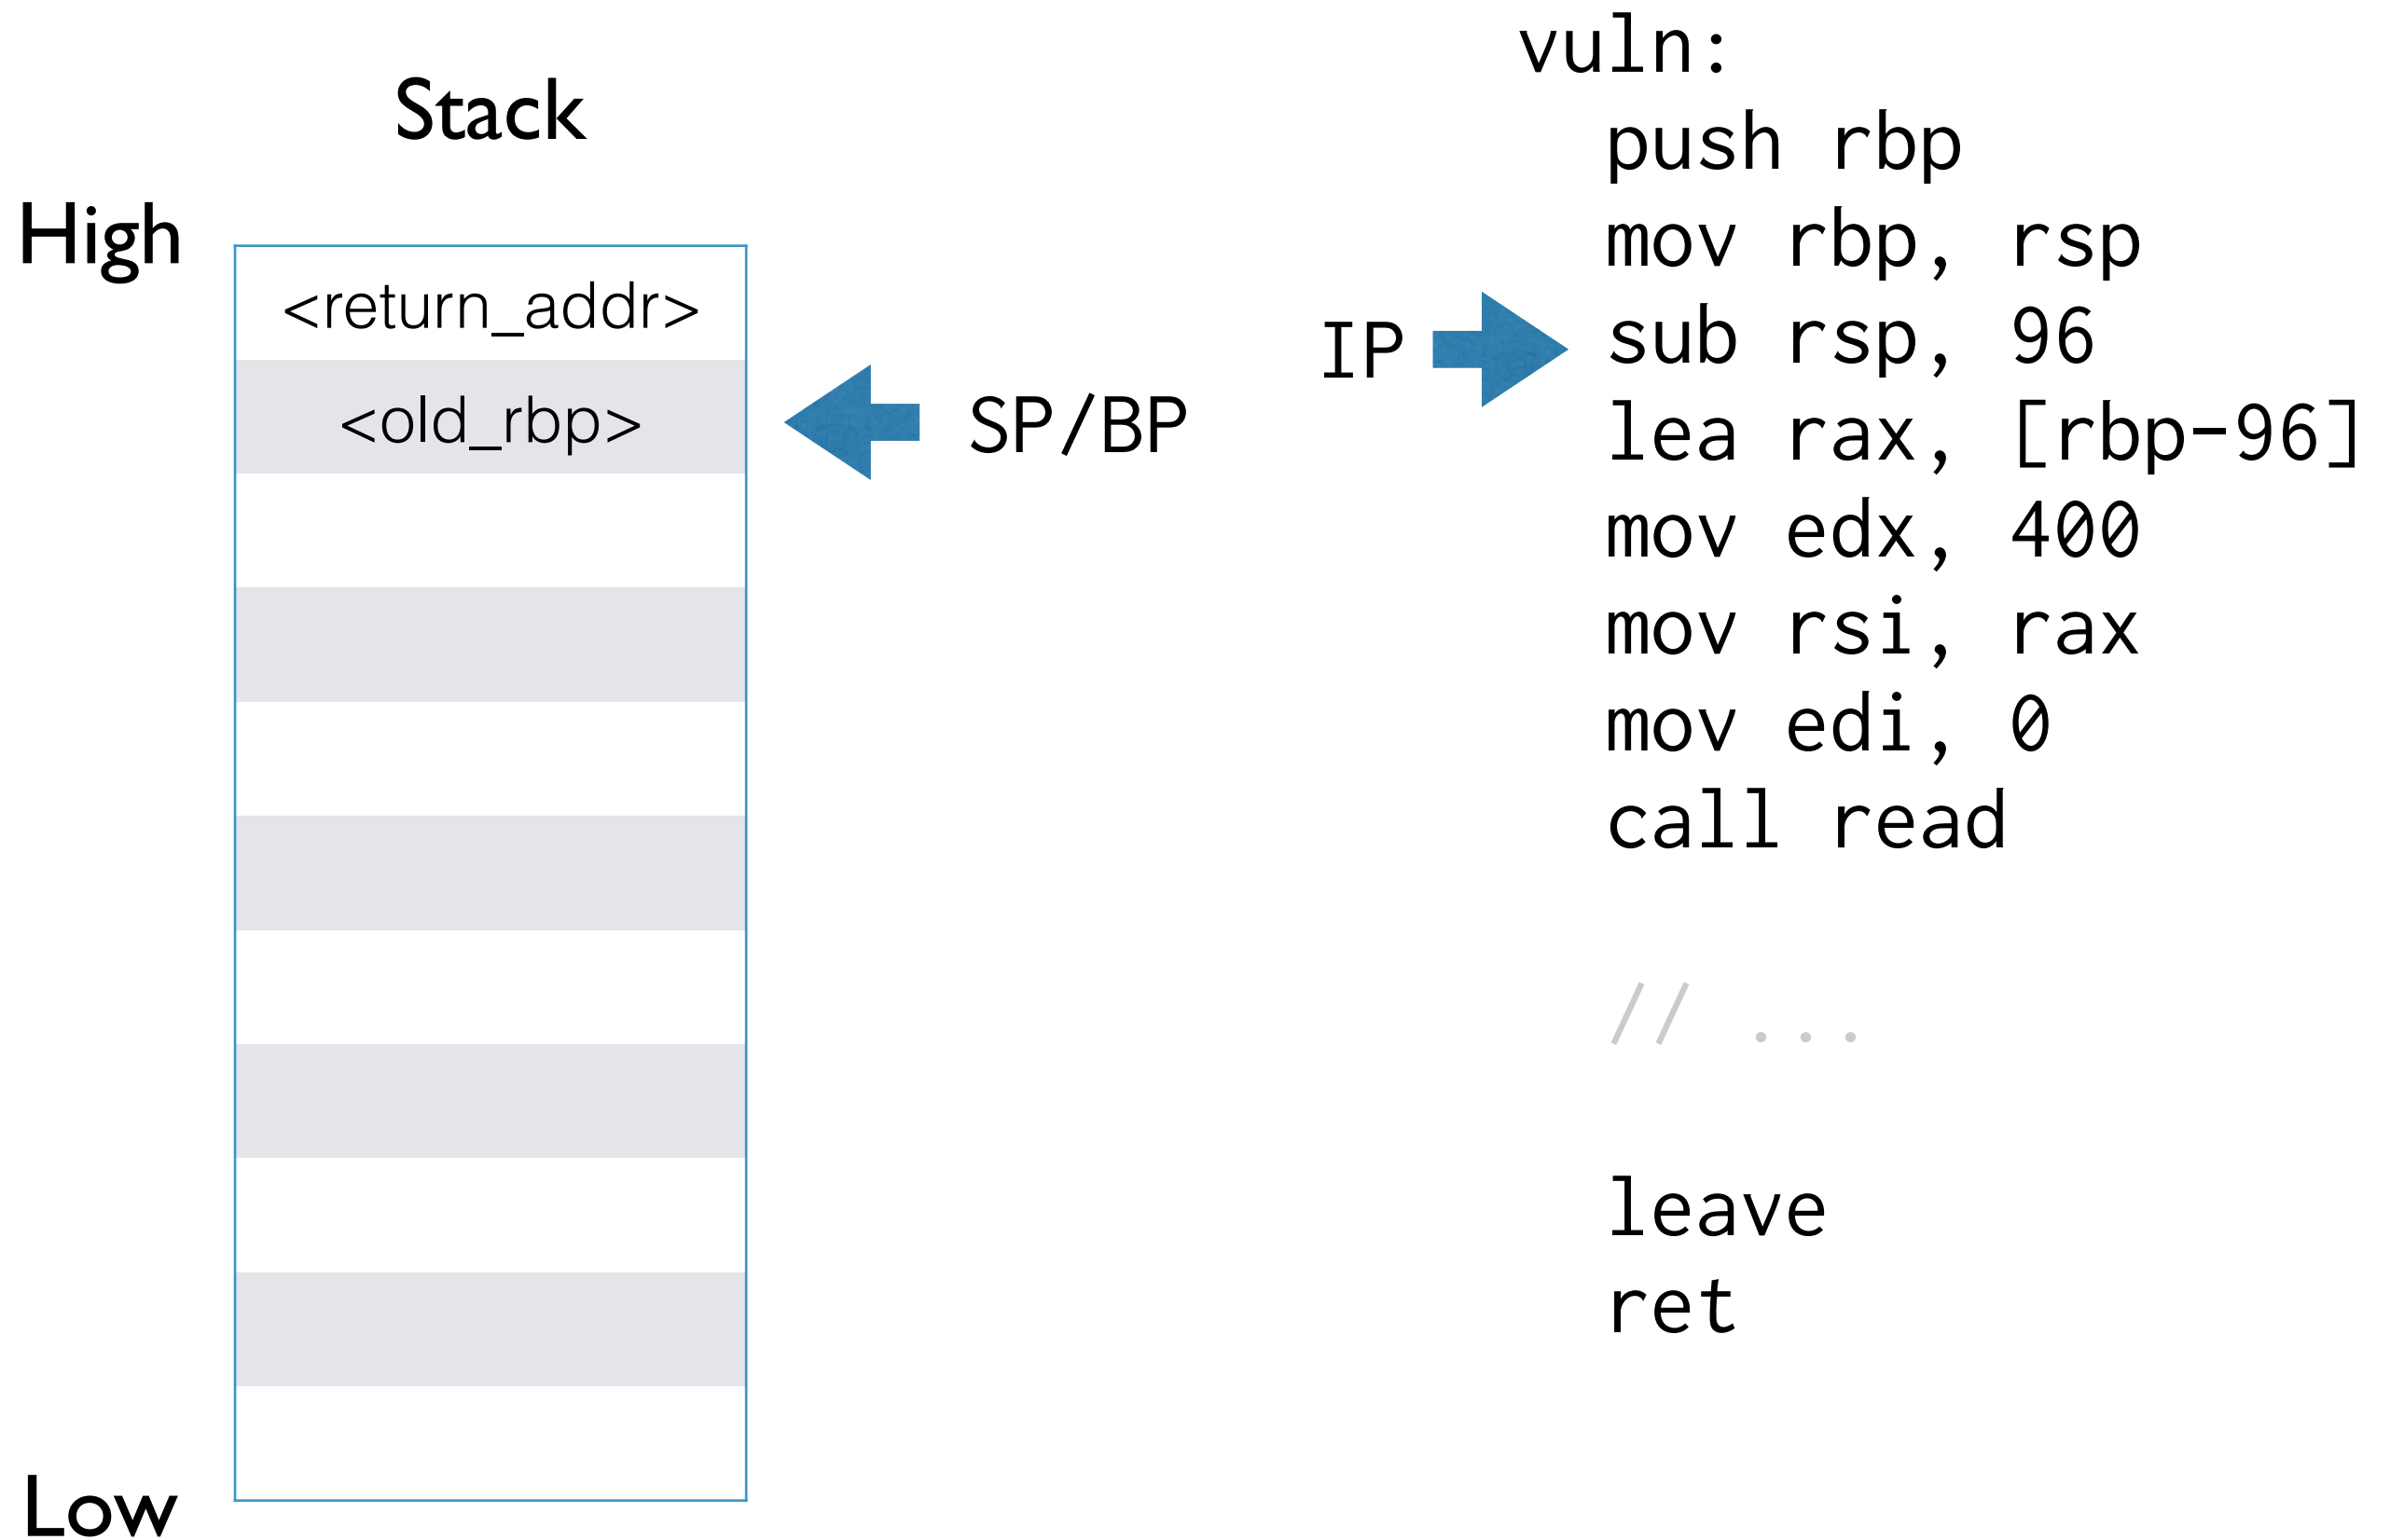
IP → vuln:

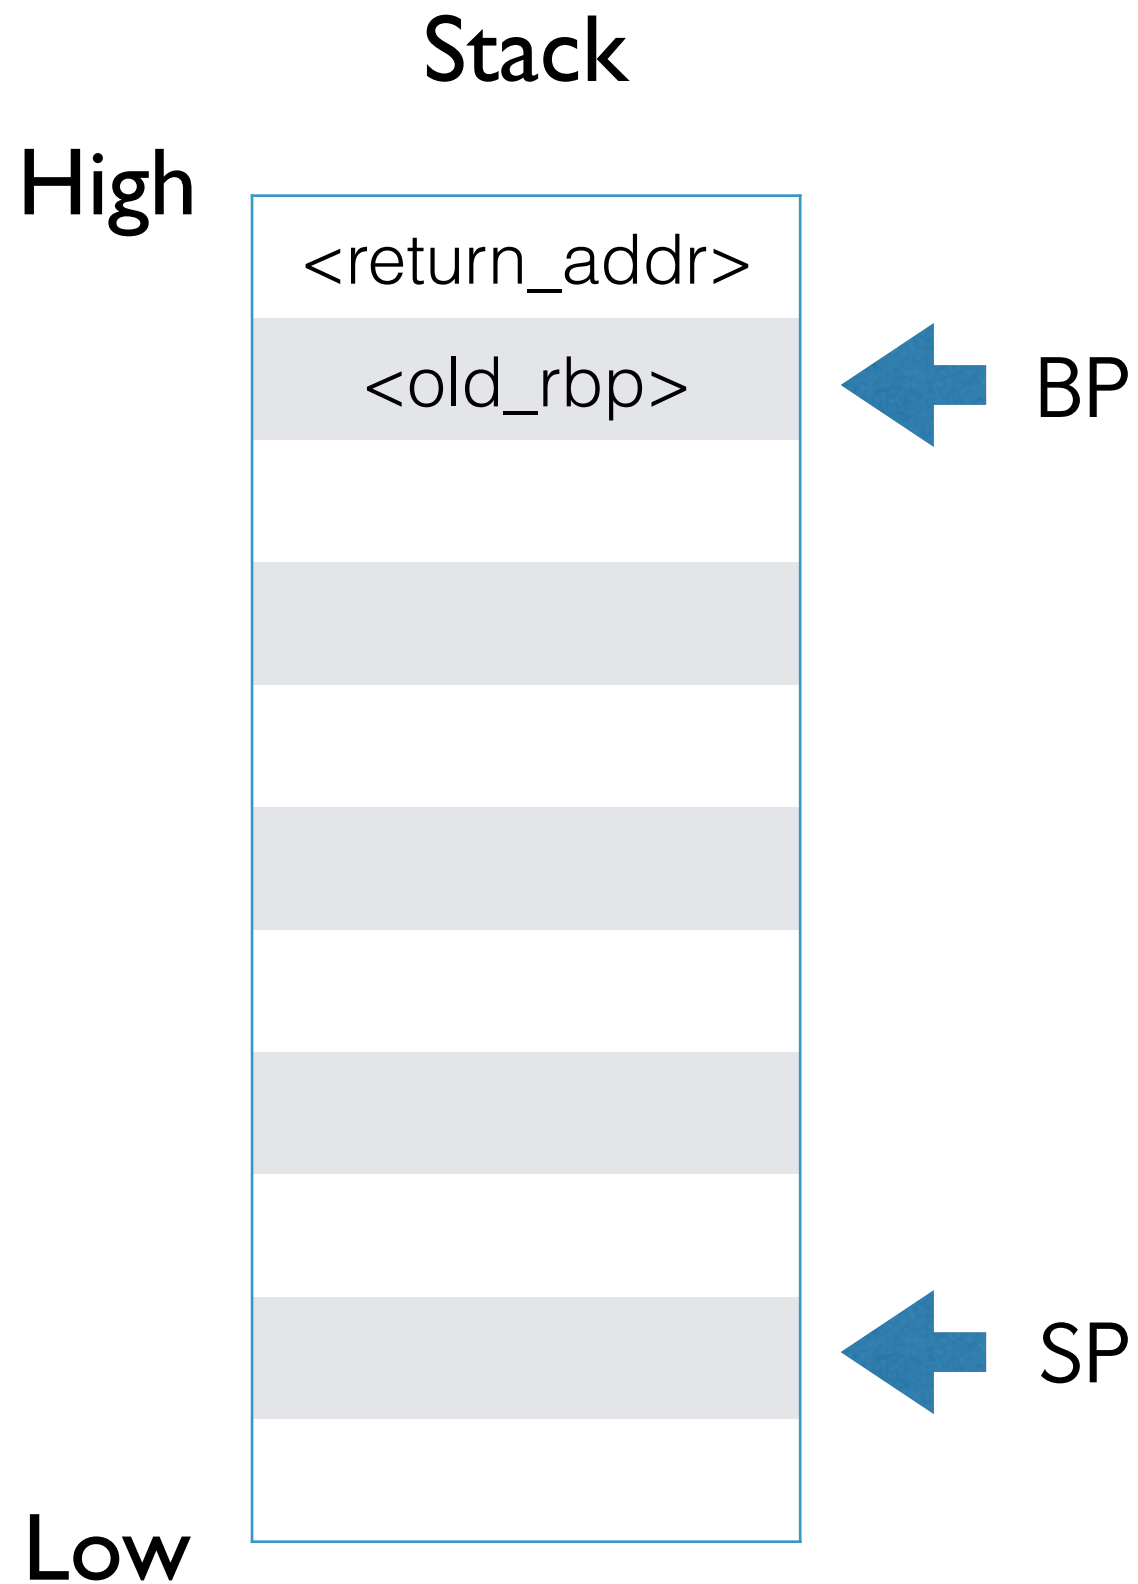
```
push rbp
mov rbp, rsp
sub rsp, 96
lea rax, [rbp-96]
mov edx, 400
mov rsi, rax
mov edi, 0
call read

// ...

leave
ret
```





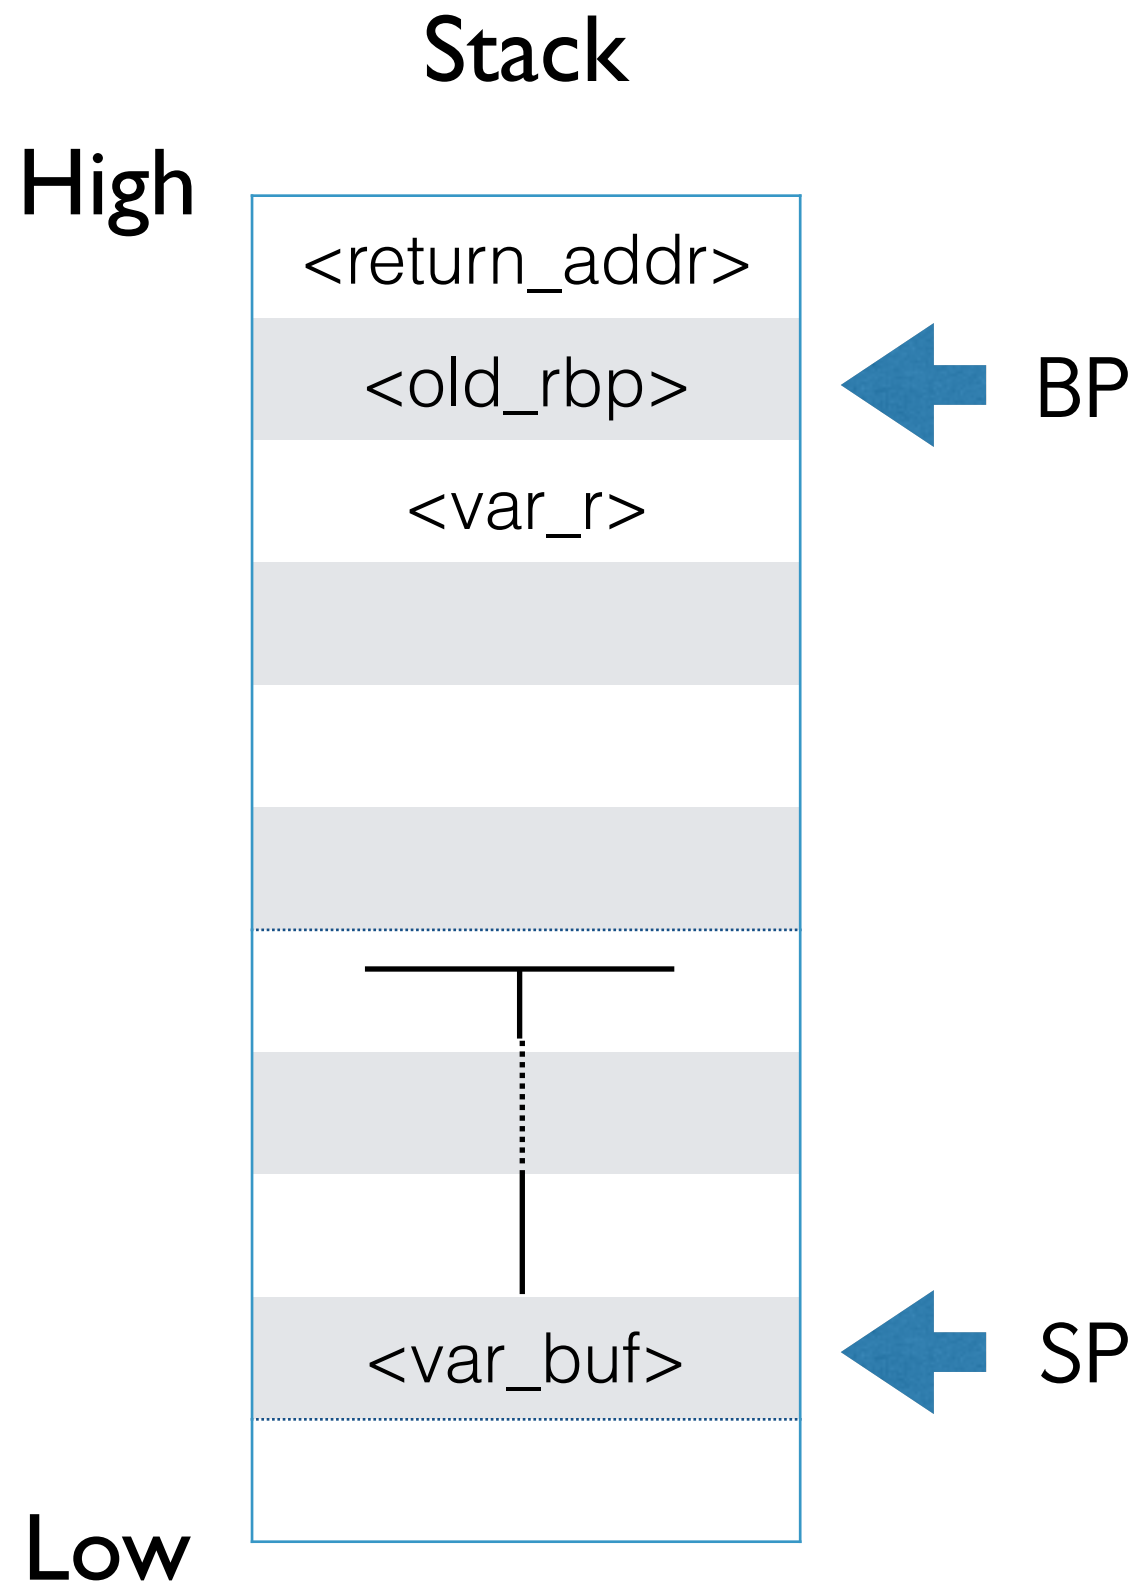


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP →



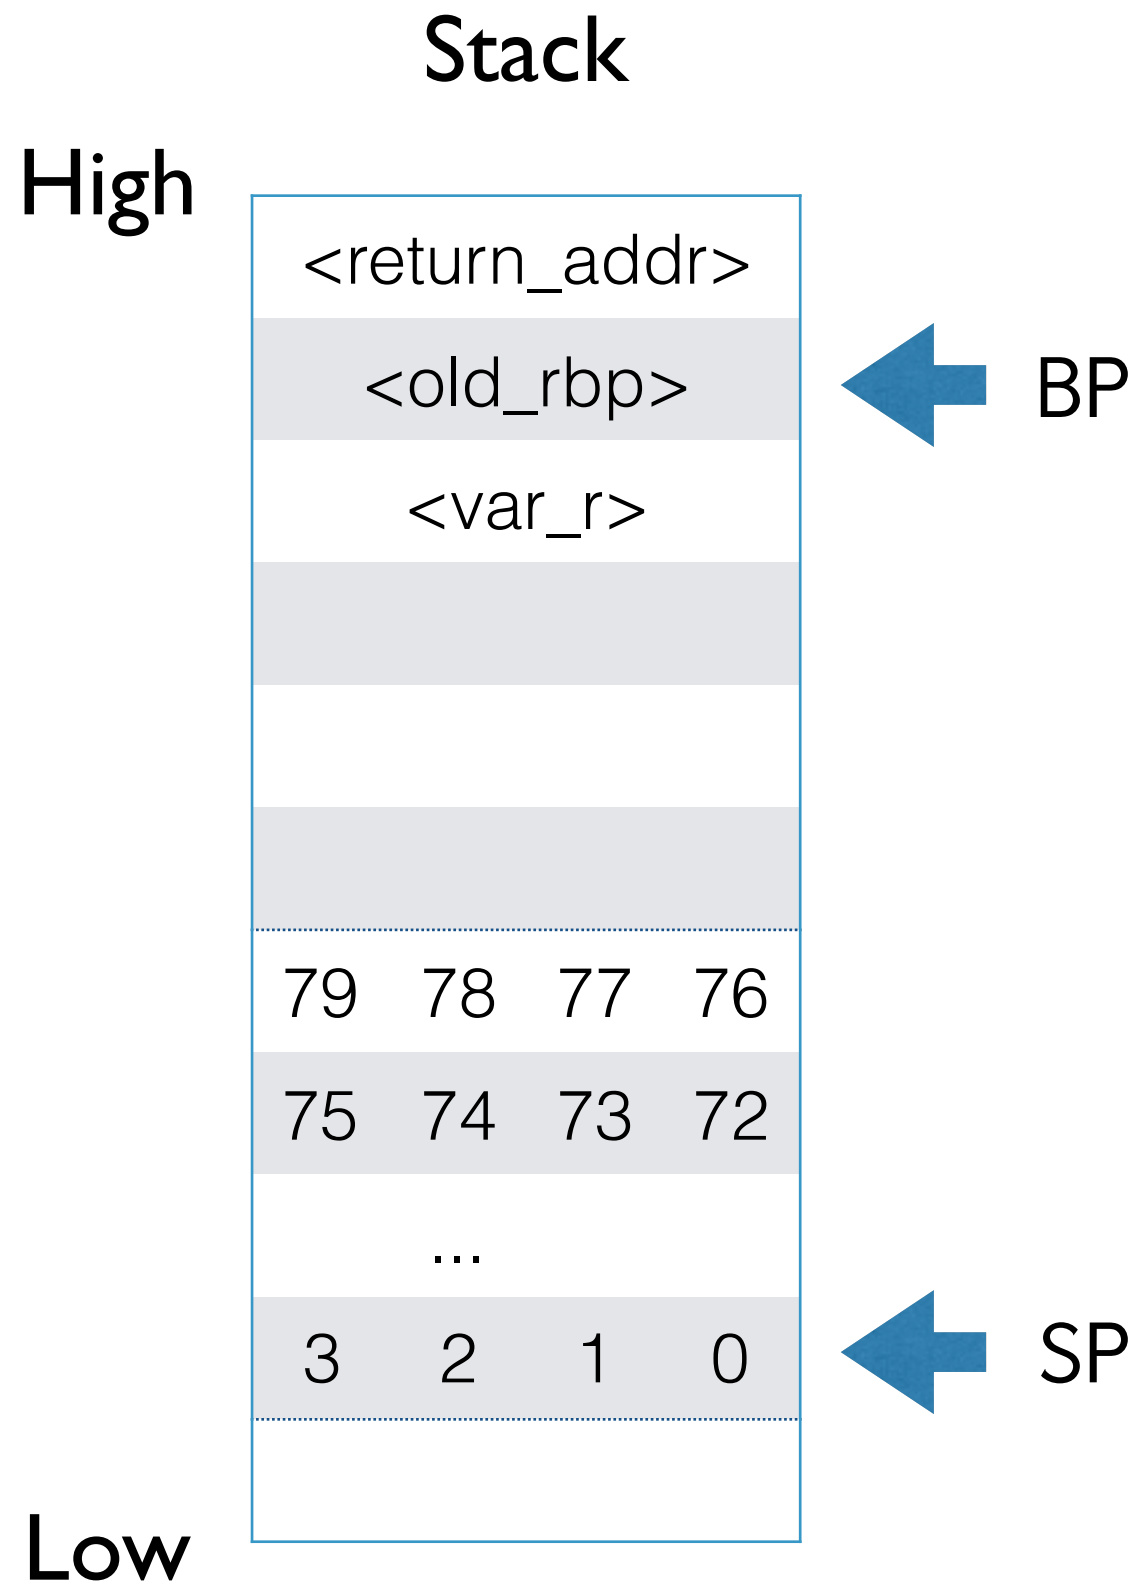
vuln:

```
push rbp
mov rbp, rsp
sub rsp, 96
lea rax, [rbp-96]
mov edx, 400
mov rsi, rax
mov edi, 0
call read

// ...

leave
ret
```

IP

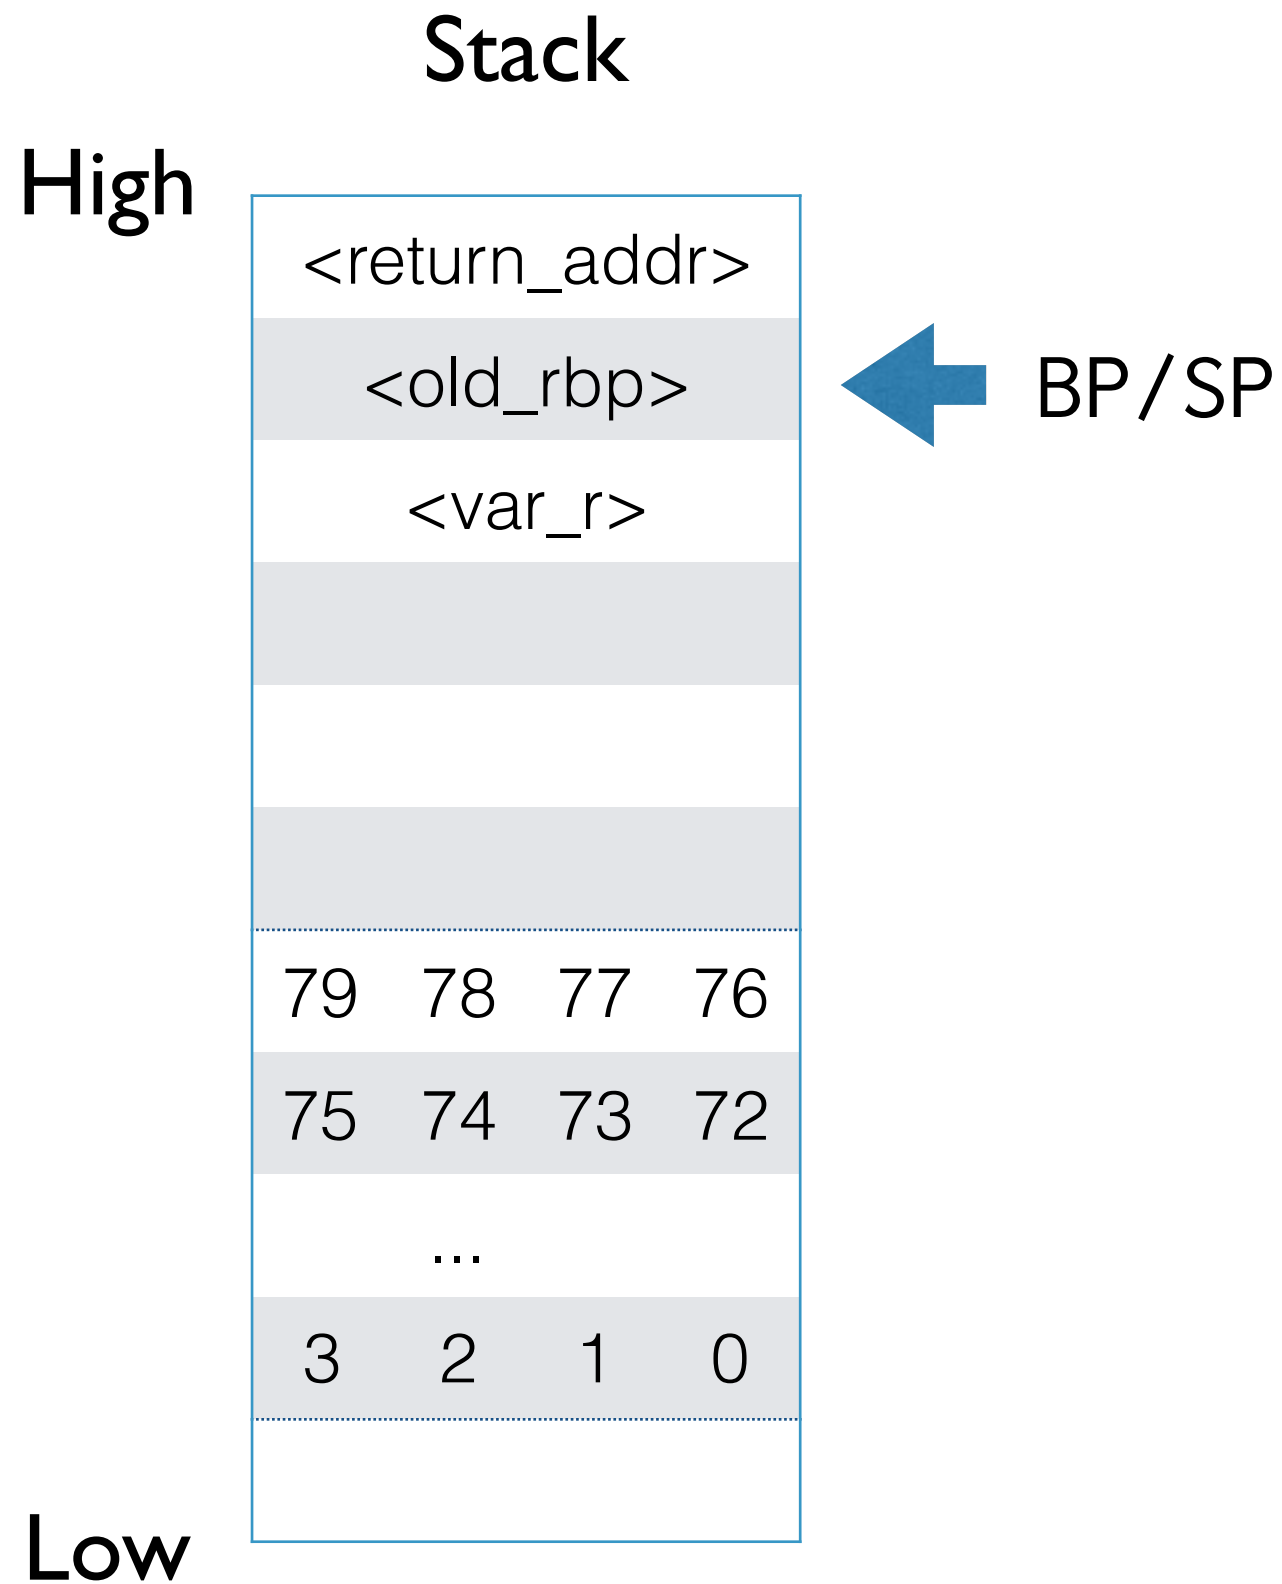


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP

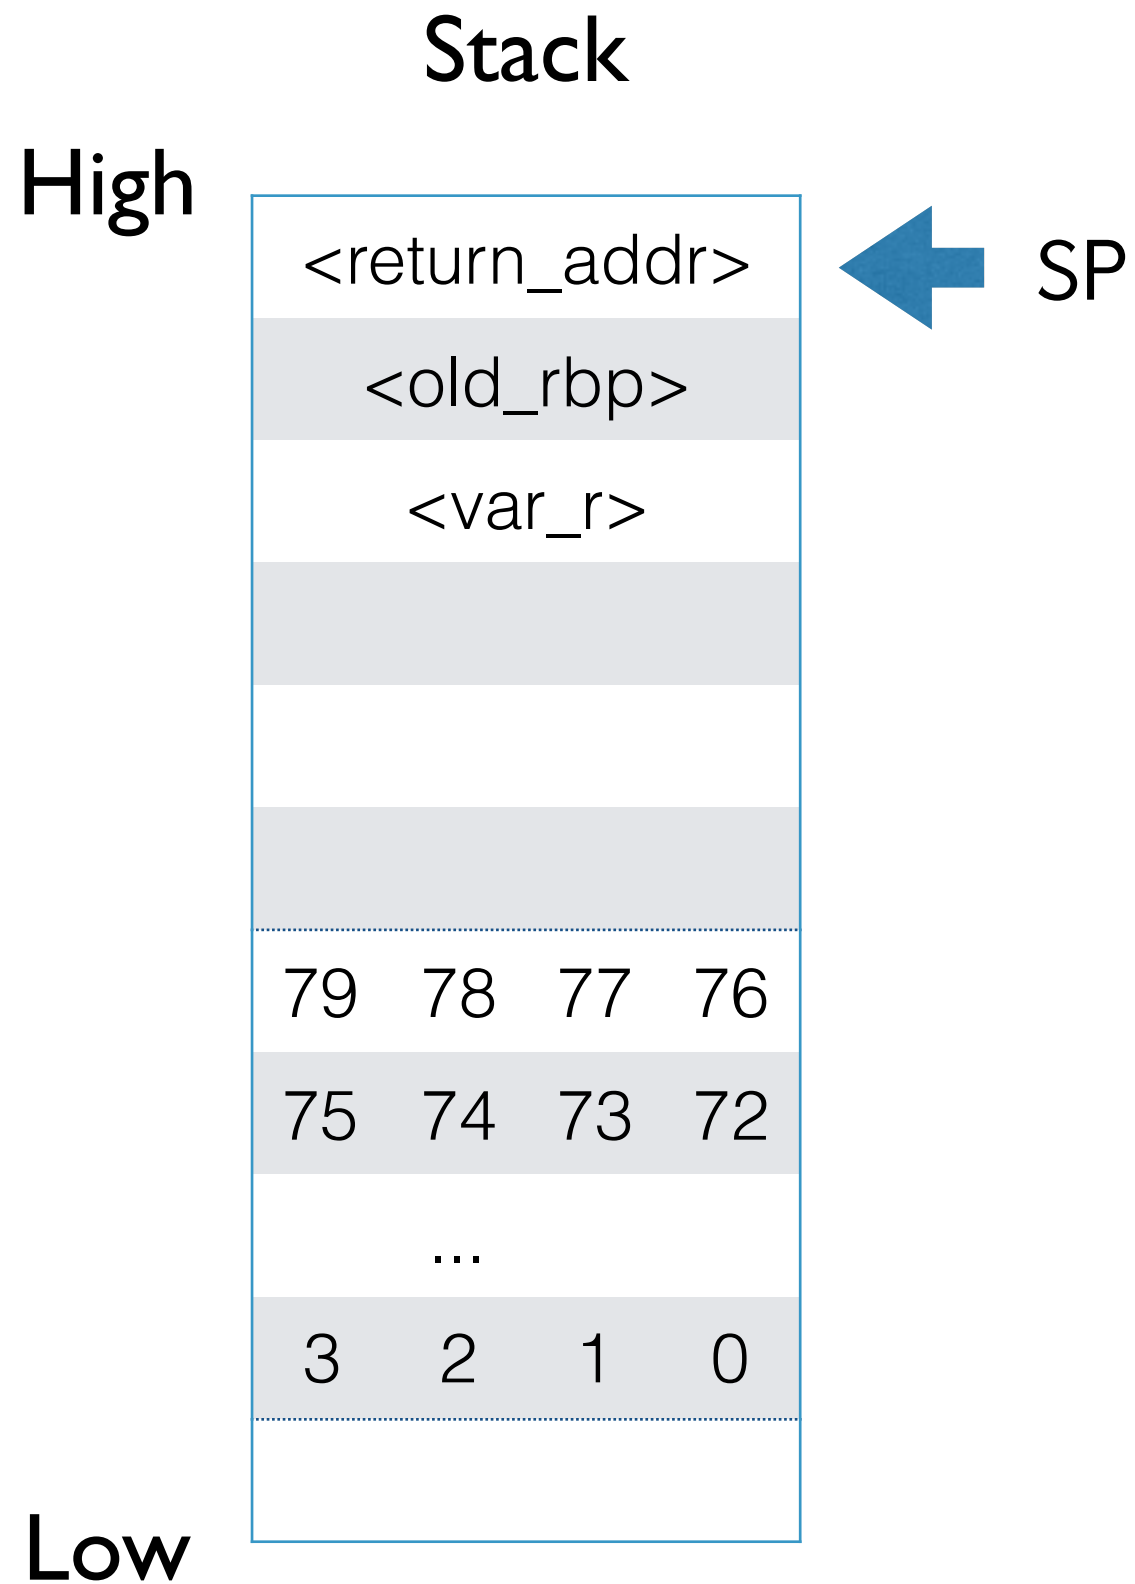


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP

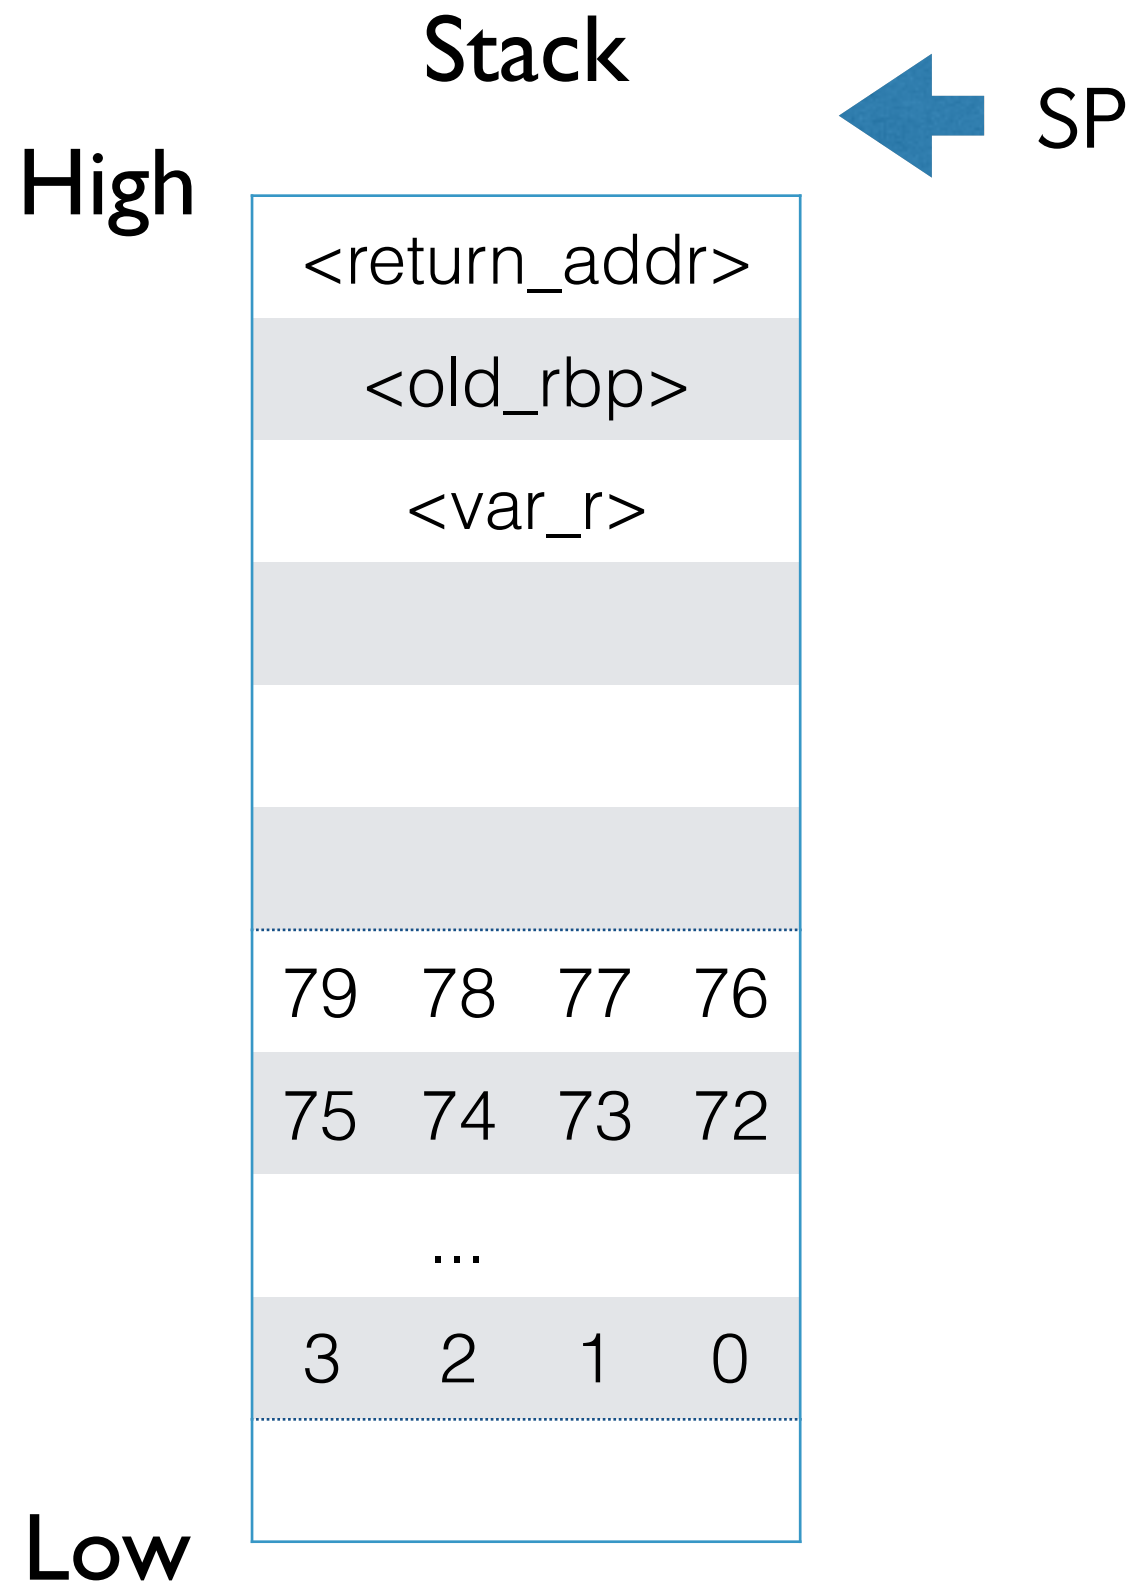


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP



```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

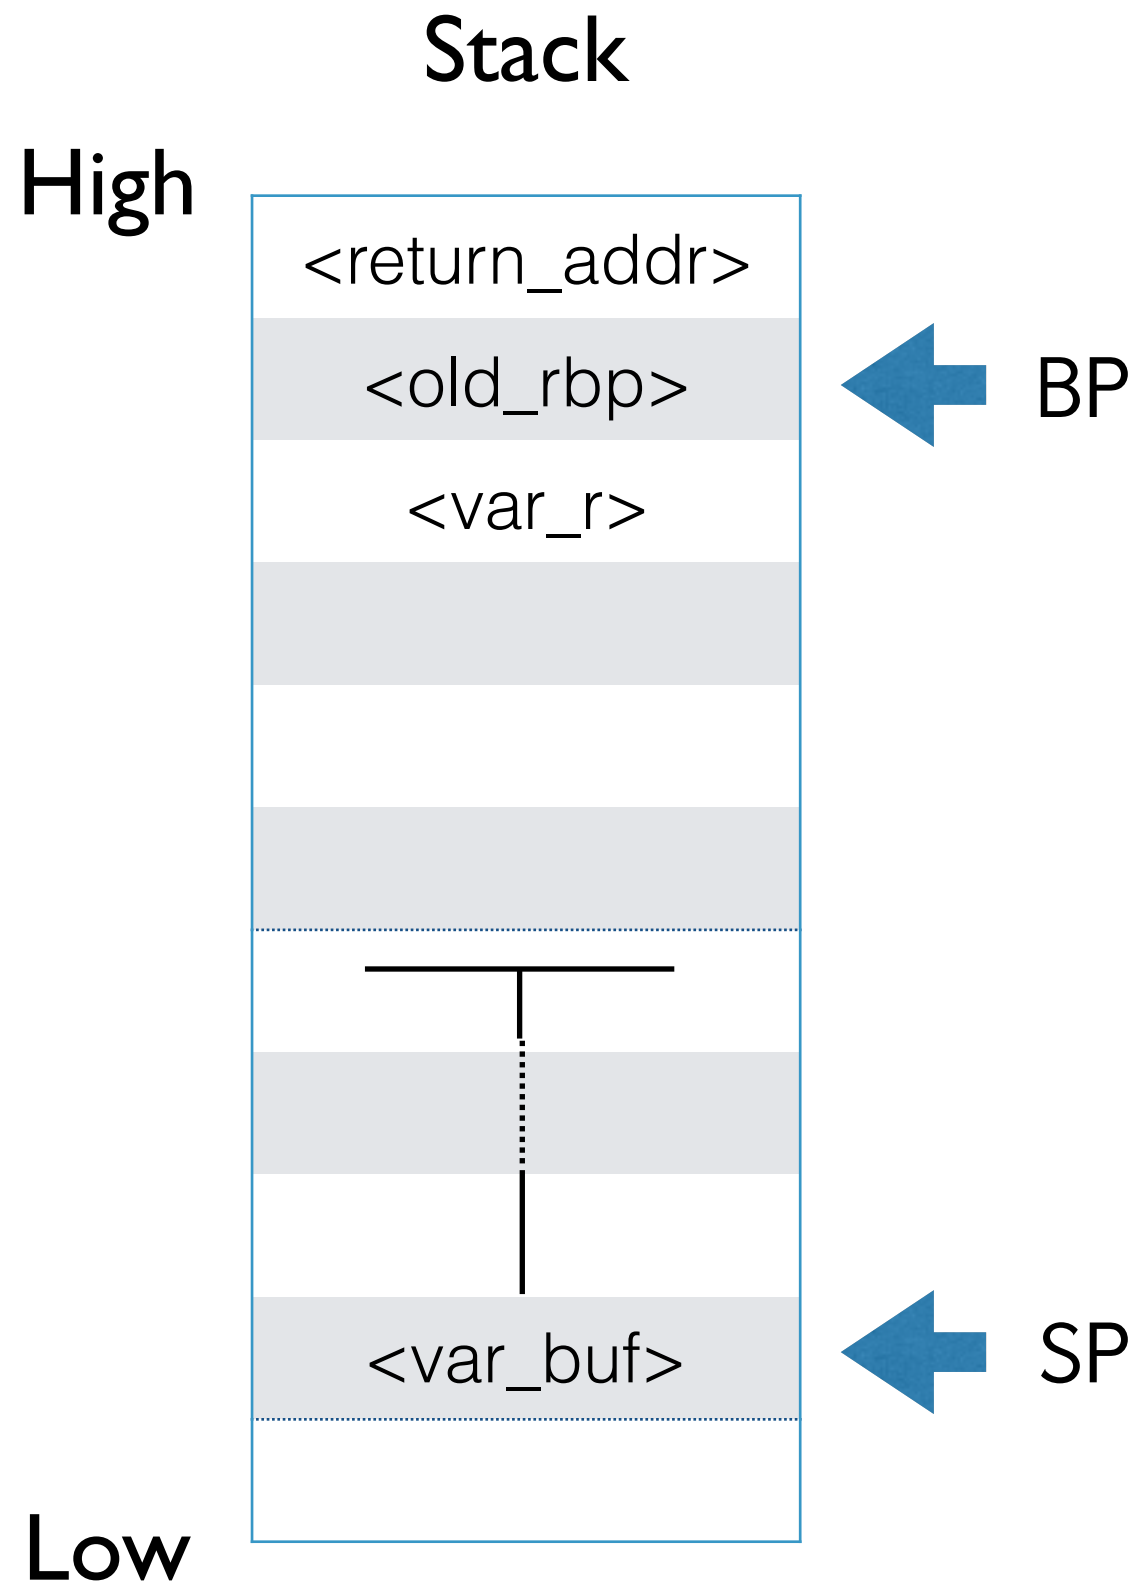
    leave
    ret
```

After function return
IP set to <return_addr>

Smashing the stack

/

Subverting control flow



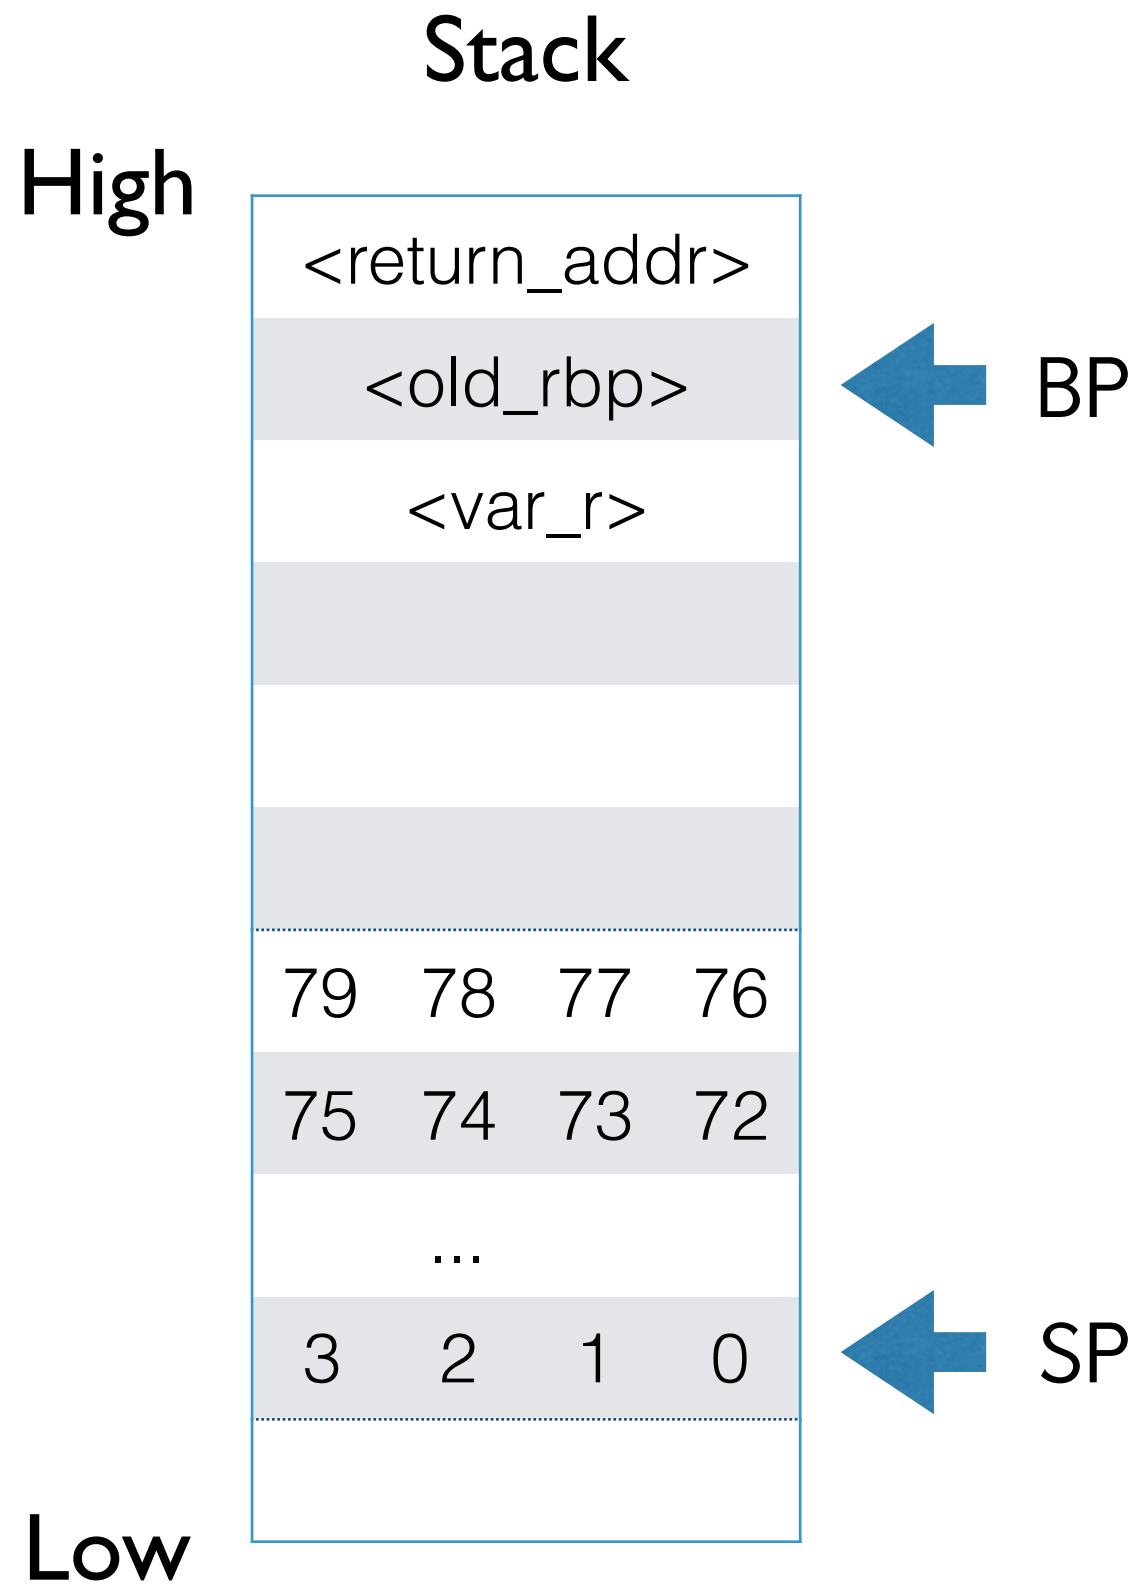
```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read
```

IP →

// ...

```
leave
ret
```

edi	rsi	edx
0	&buf	400

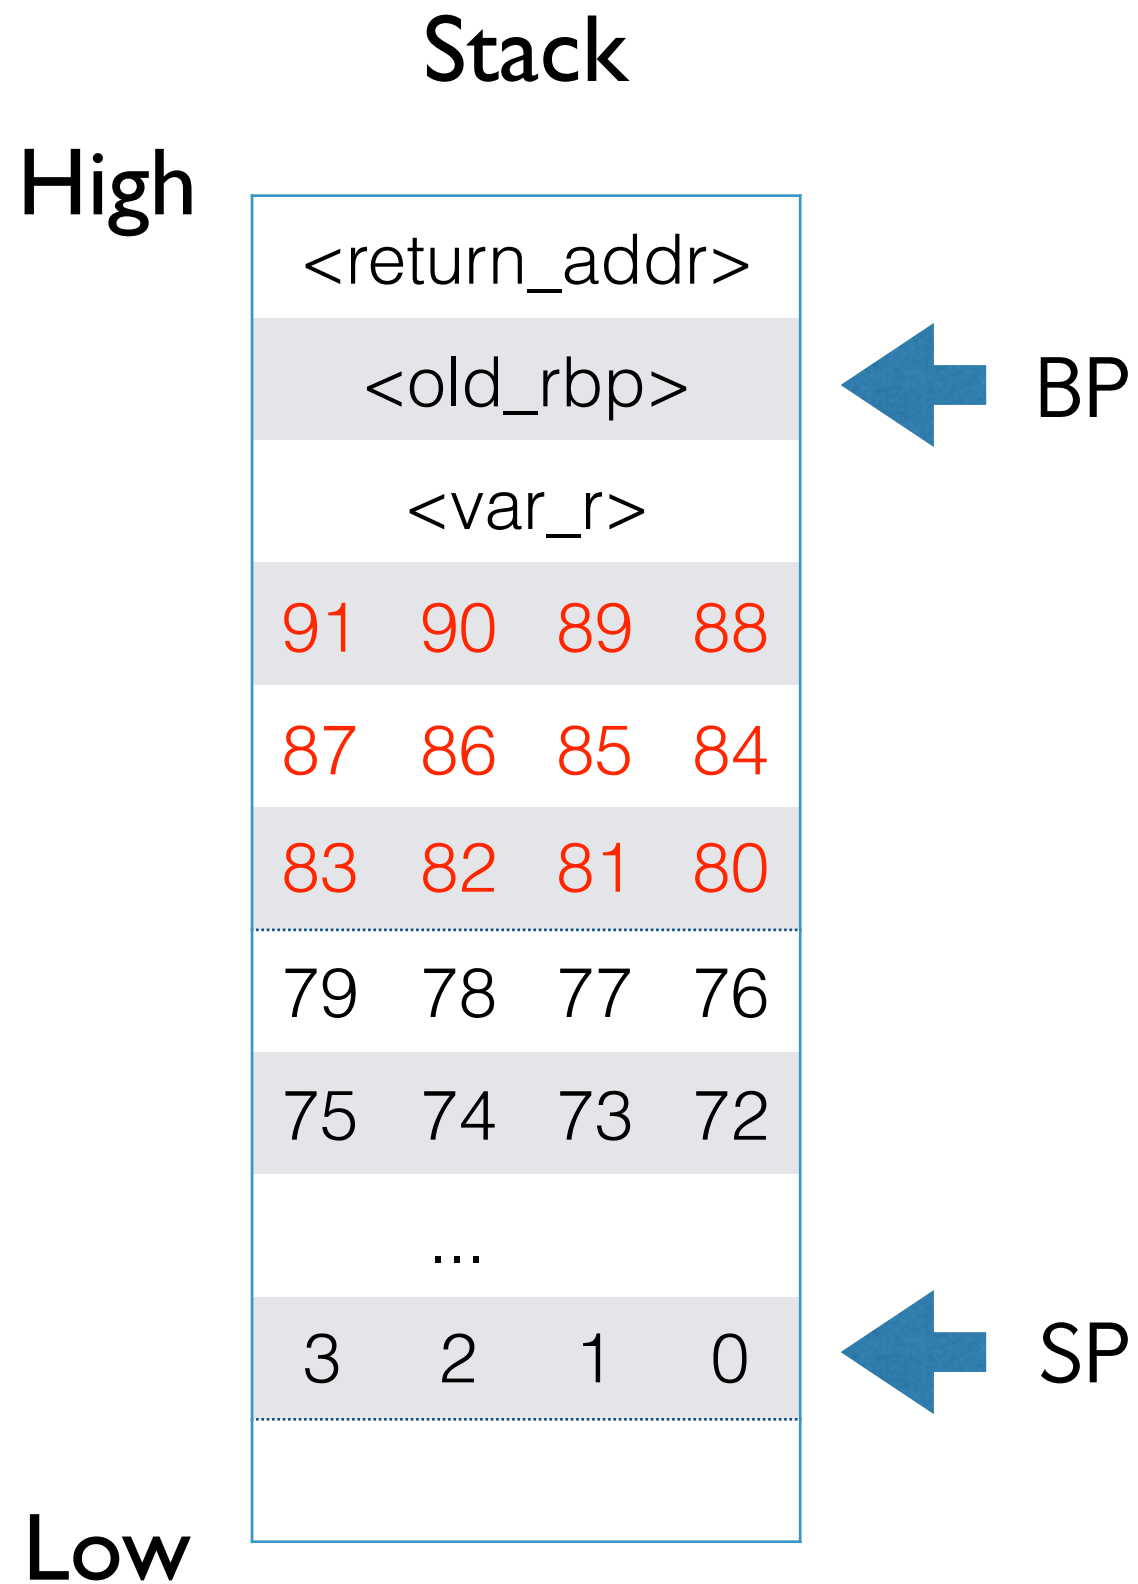


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP

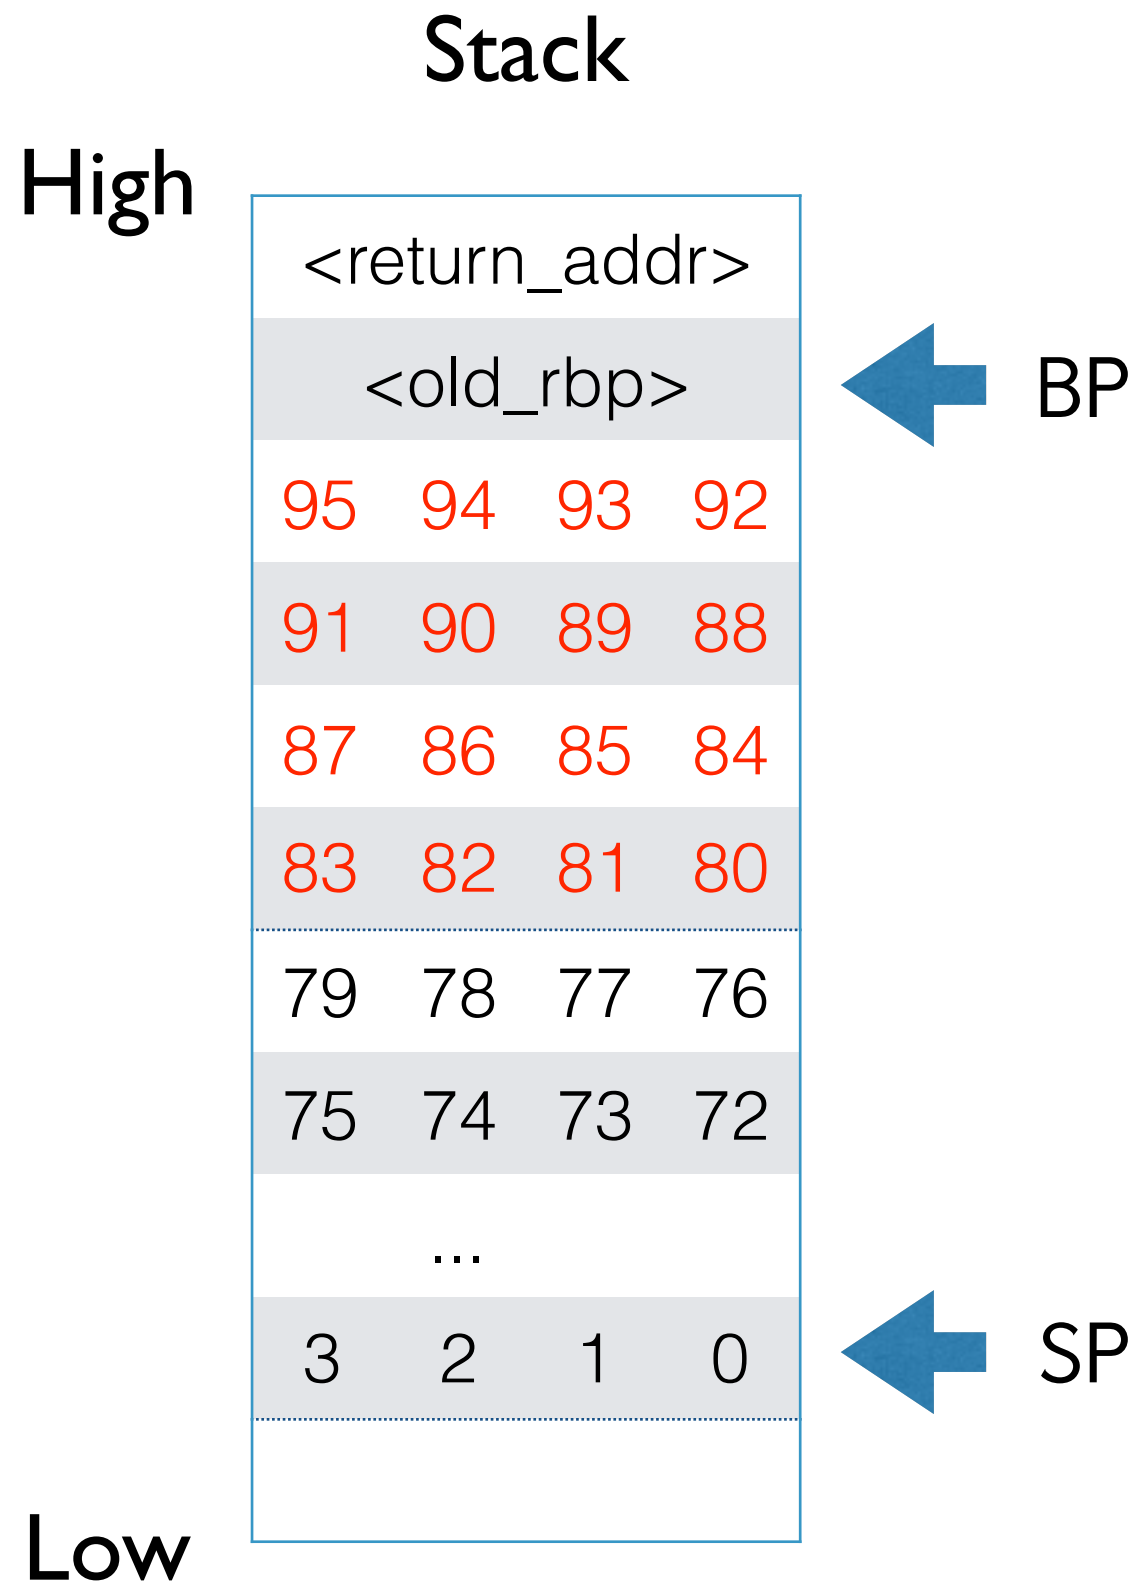


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP

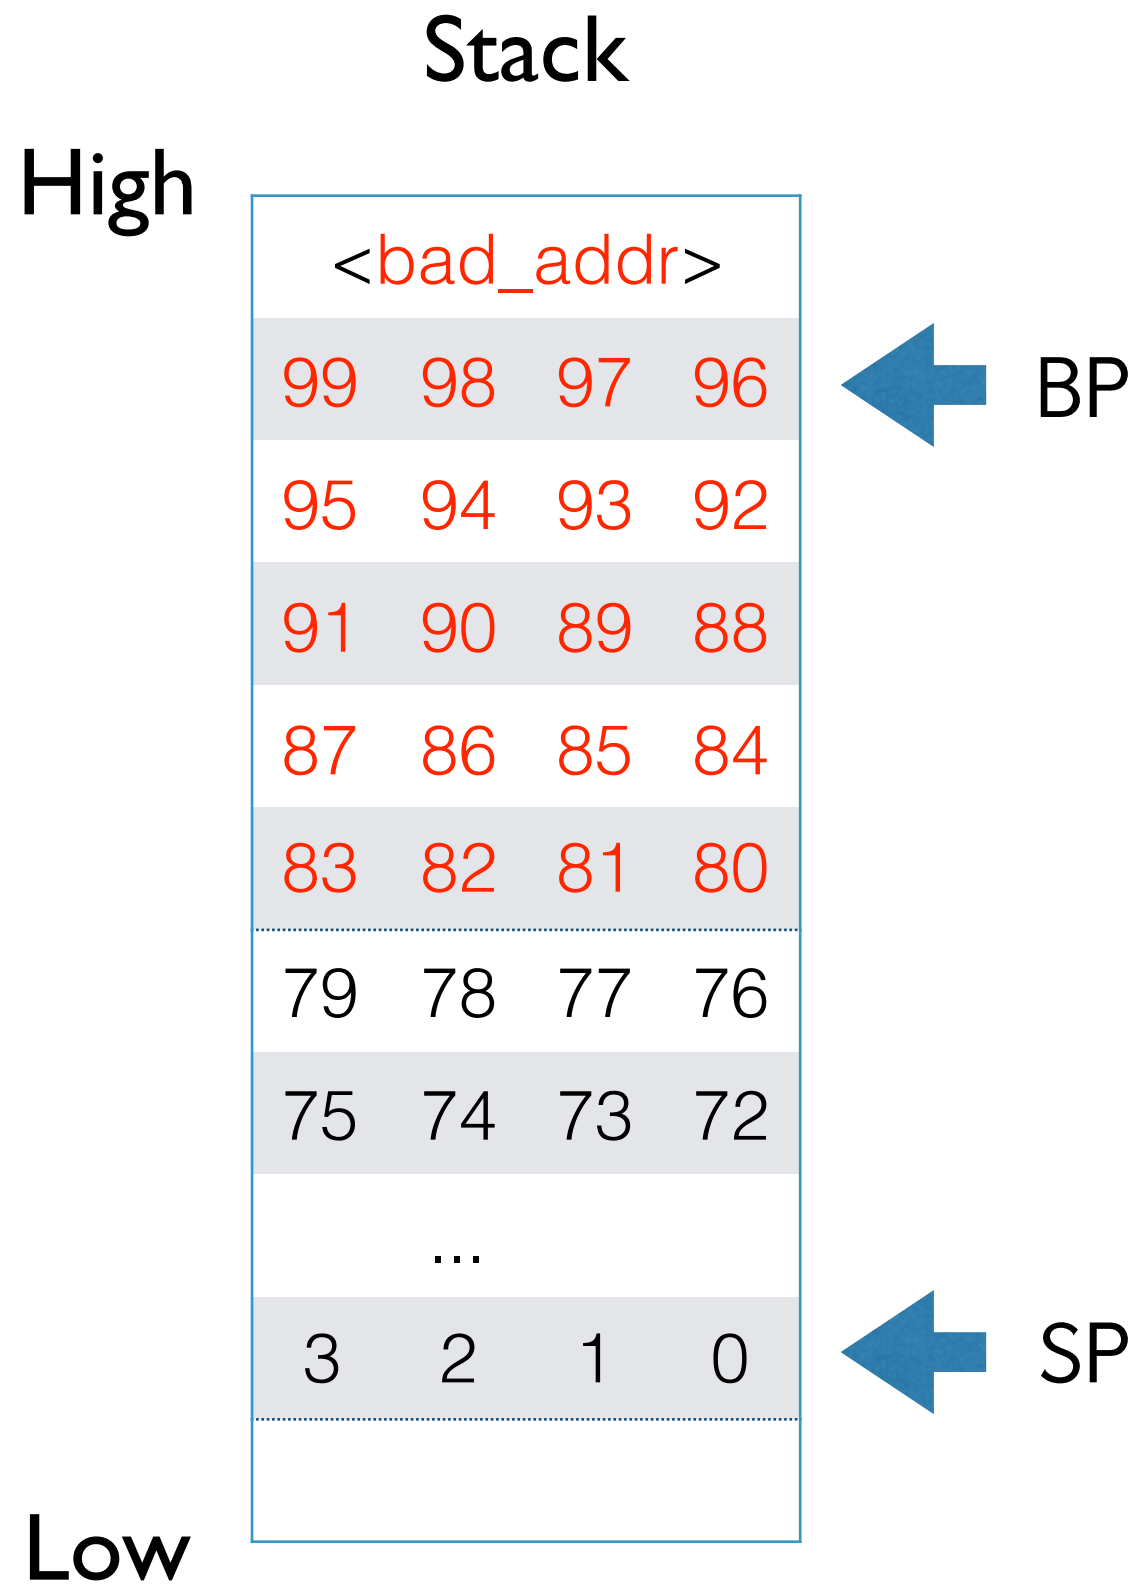


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP

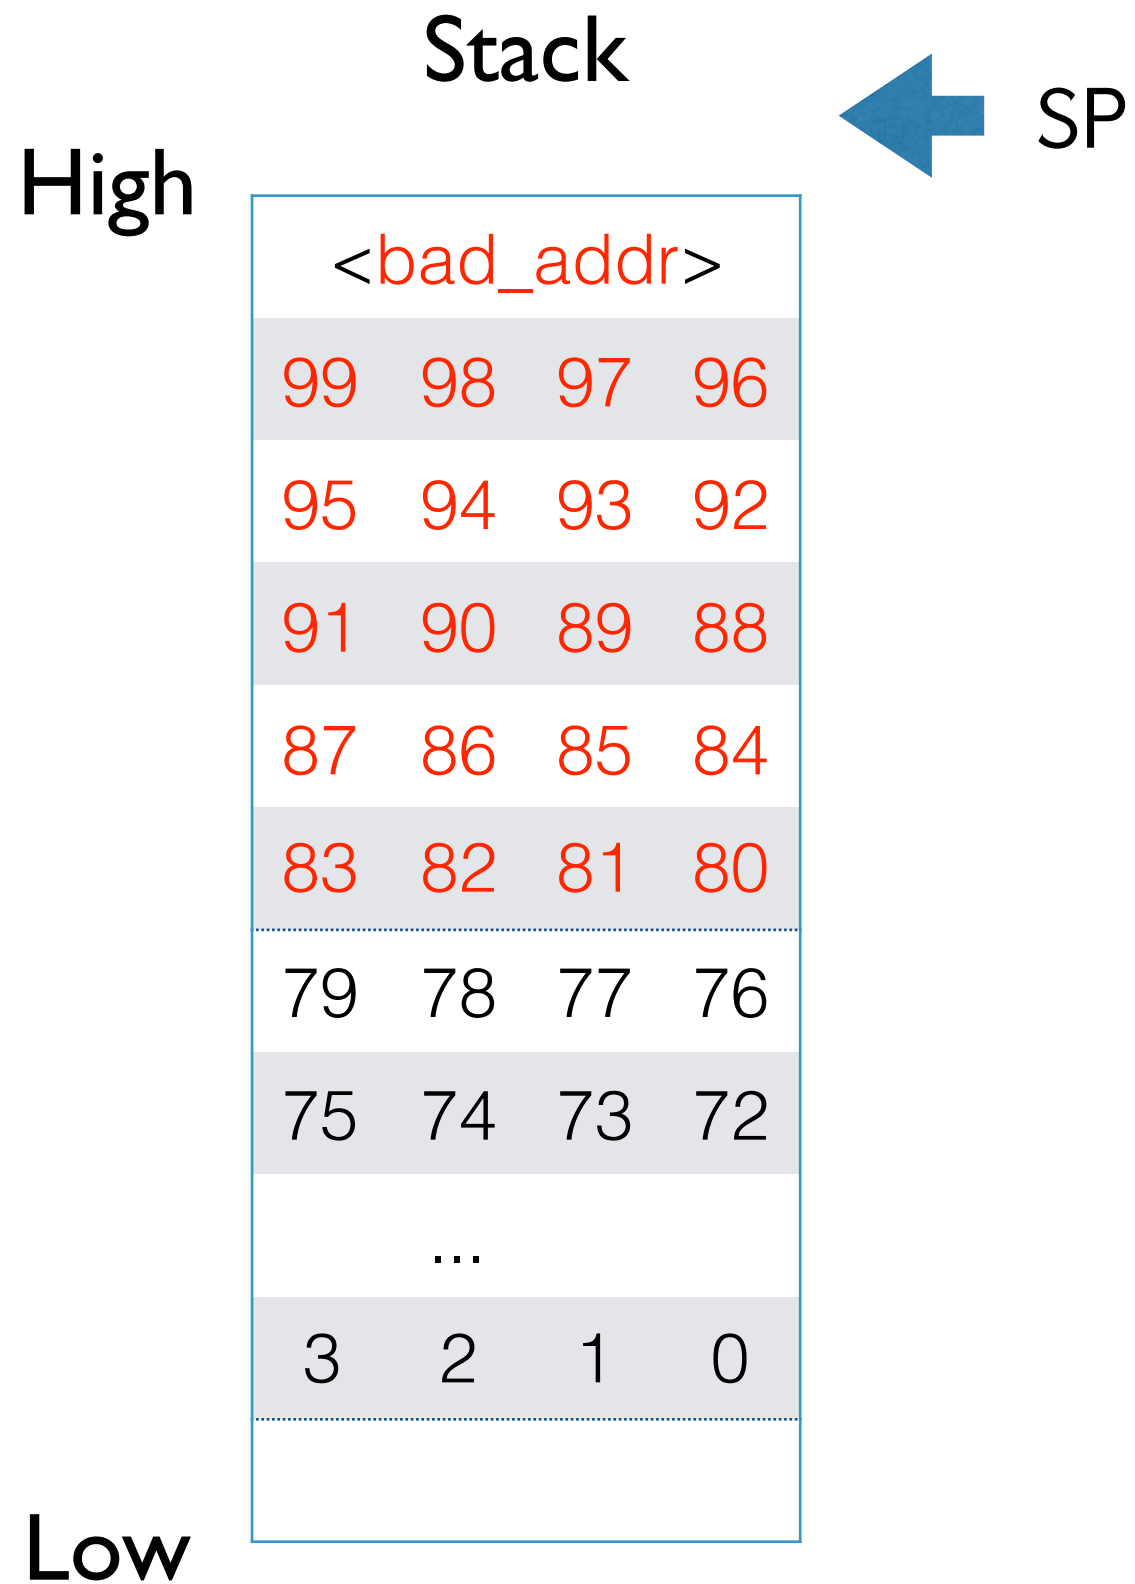


```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

    leave
    ret
```

IP



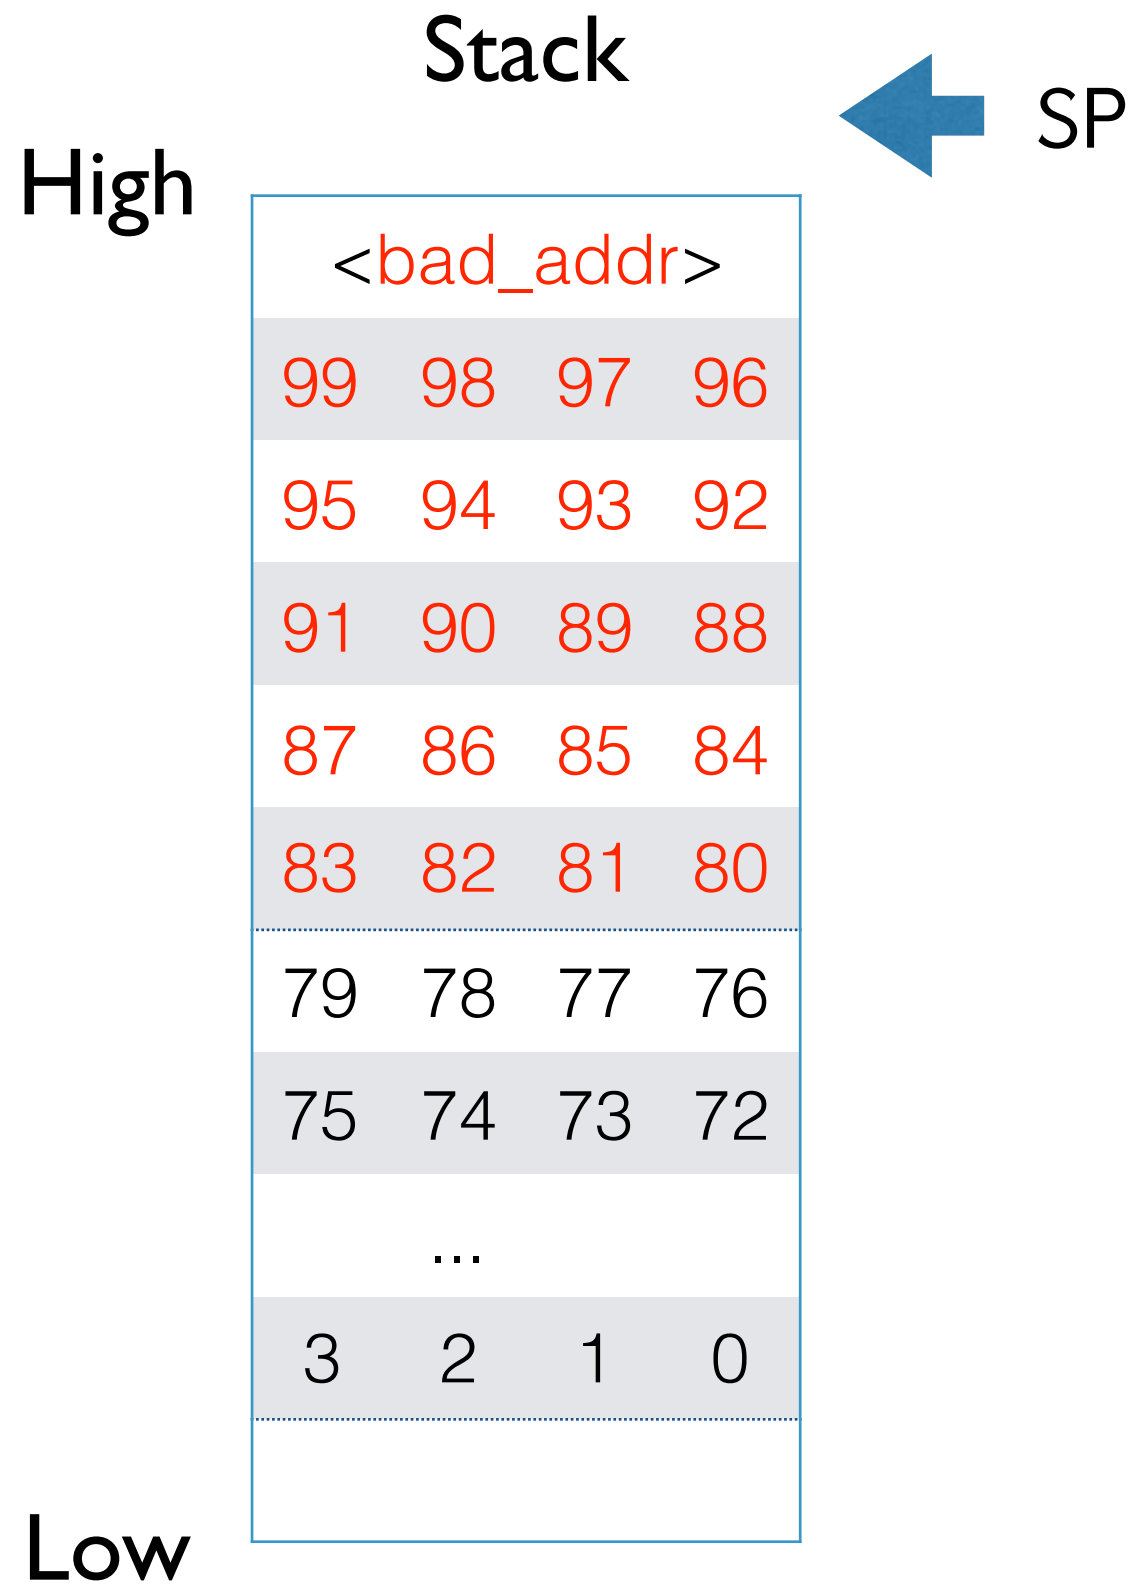
vuln:

```
push rbp
mov rbp, rsp
sub rsp, 96
lea rax, [rbp-96]
mov edx, 400
mov rsi, rax
mov edi, 0
call read
```

// ...

leave

IP → ret



```
vuln:
    push rbp
    mov rbp, rsp
    sub rsp, 96
    lea rax, [rbp-96]
    mov edx, 400
    mov rsi, rax
    mov edi, 0
    call read

    // ...

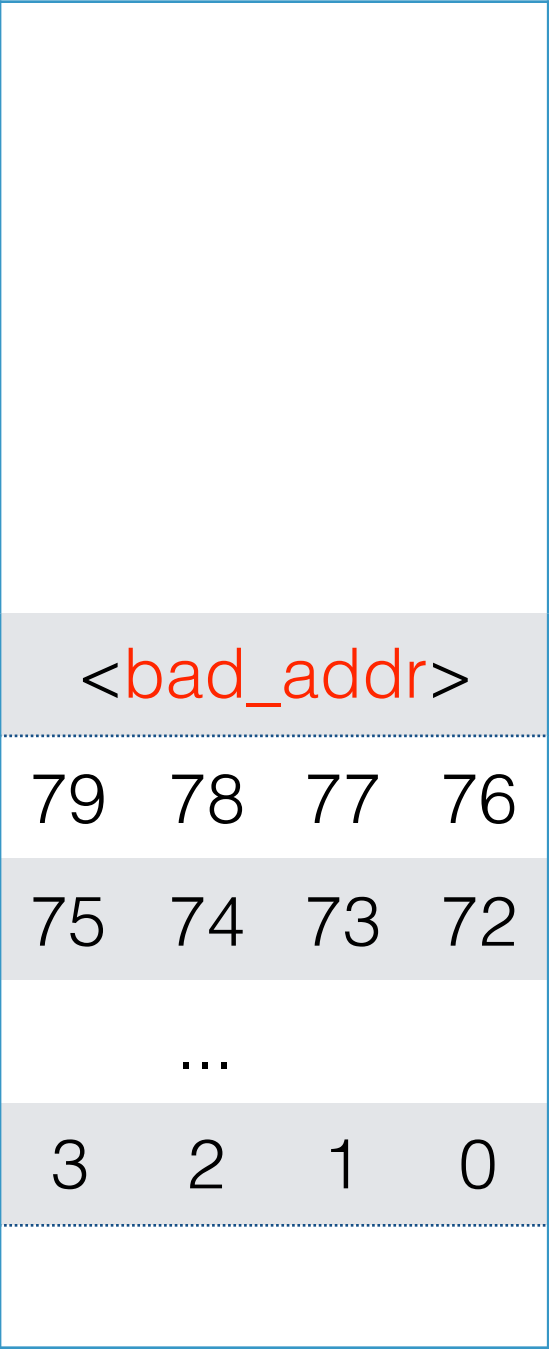
    leave
    ret
```

After function return
IP set to <bad_addr>

Where do we jump to?

Stack

High



Low

Stack

High

payload

<bad_addr>

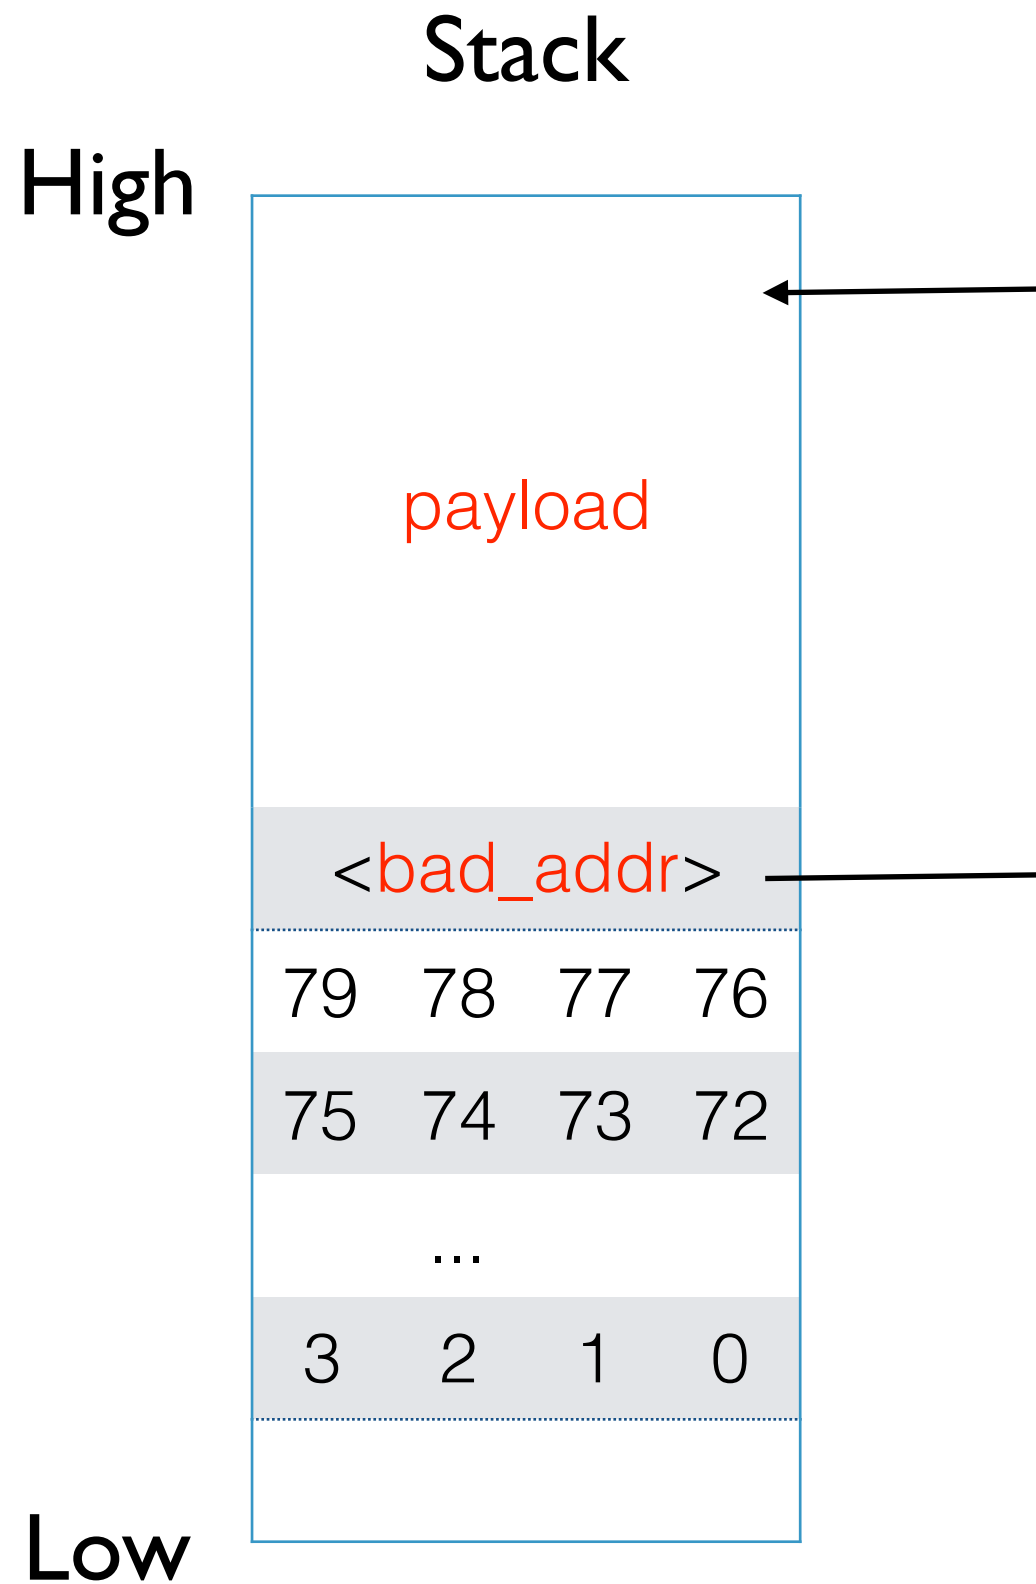
79 78 77 76

75 74 73 72

...

3 2 1 0

Low



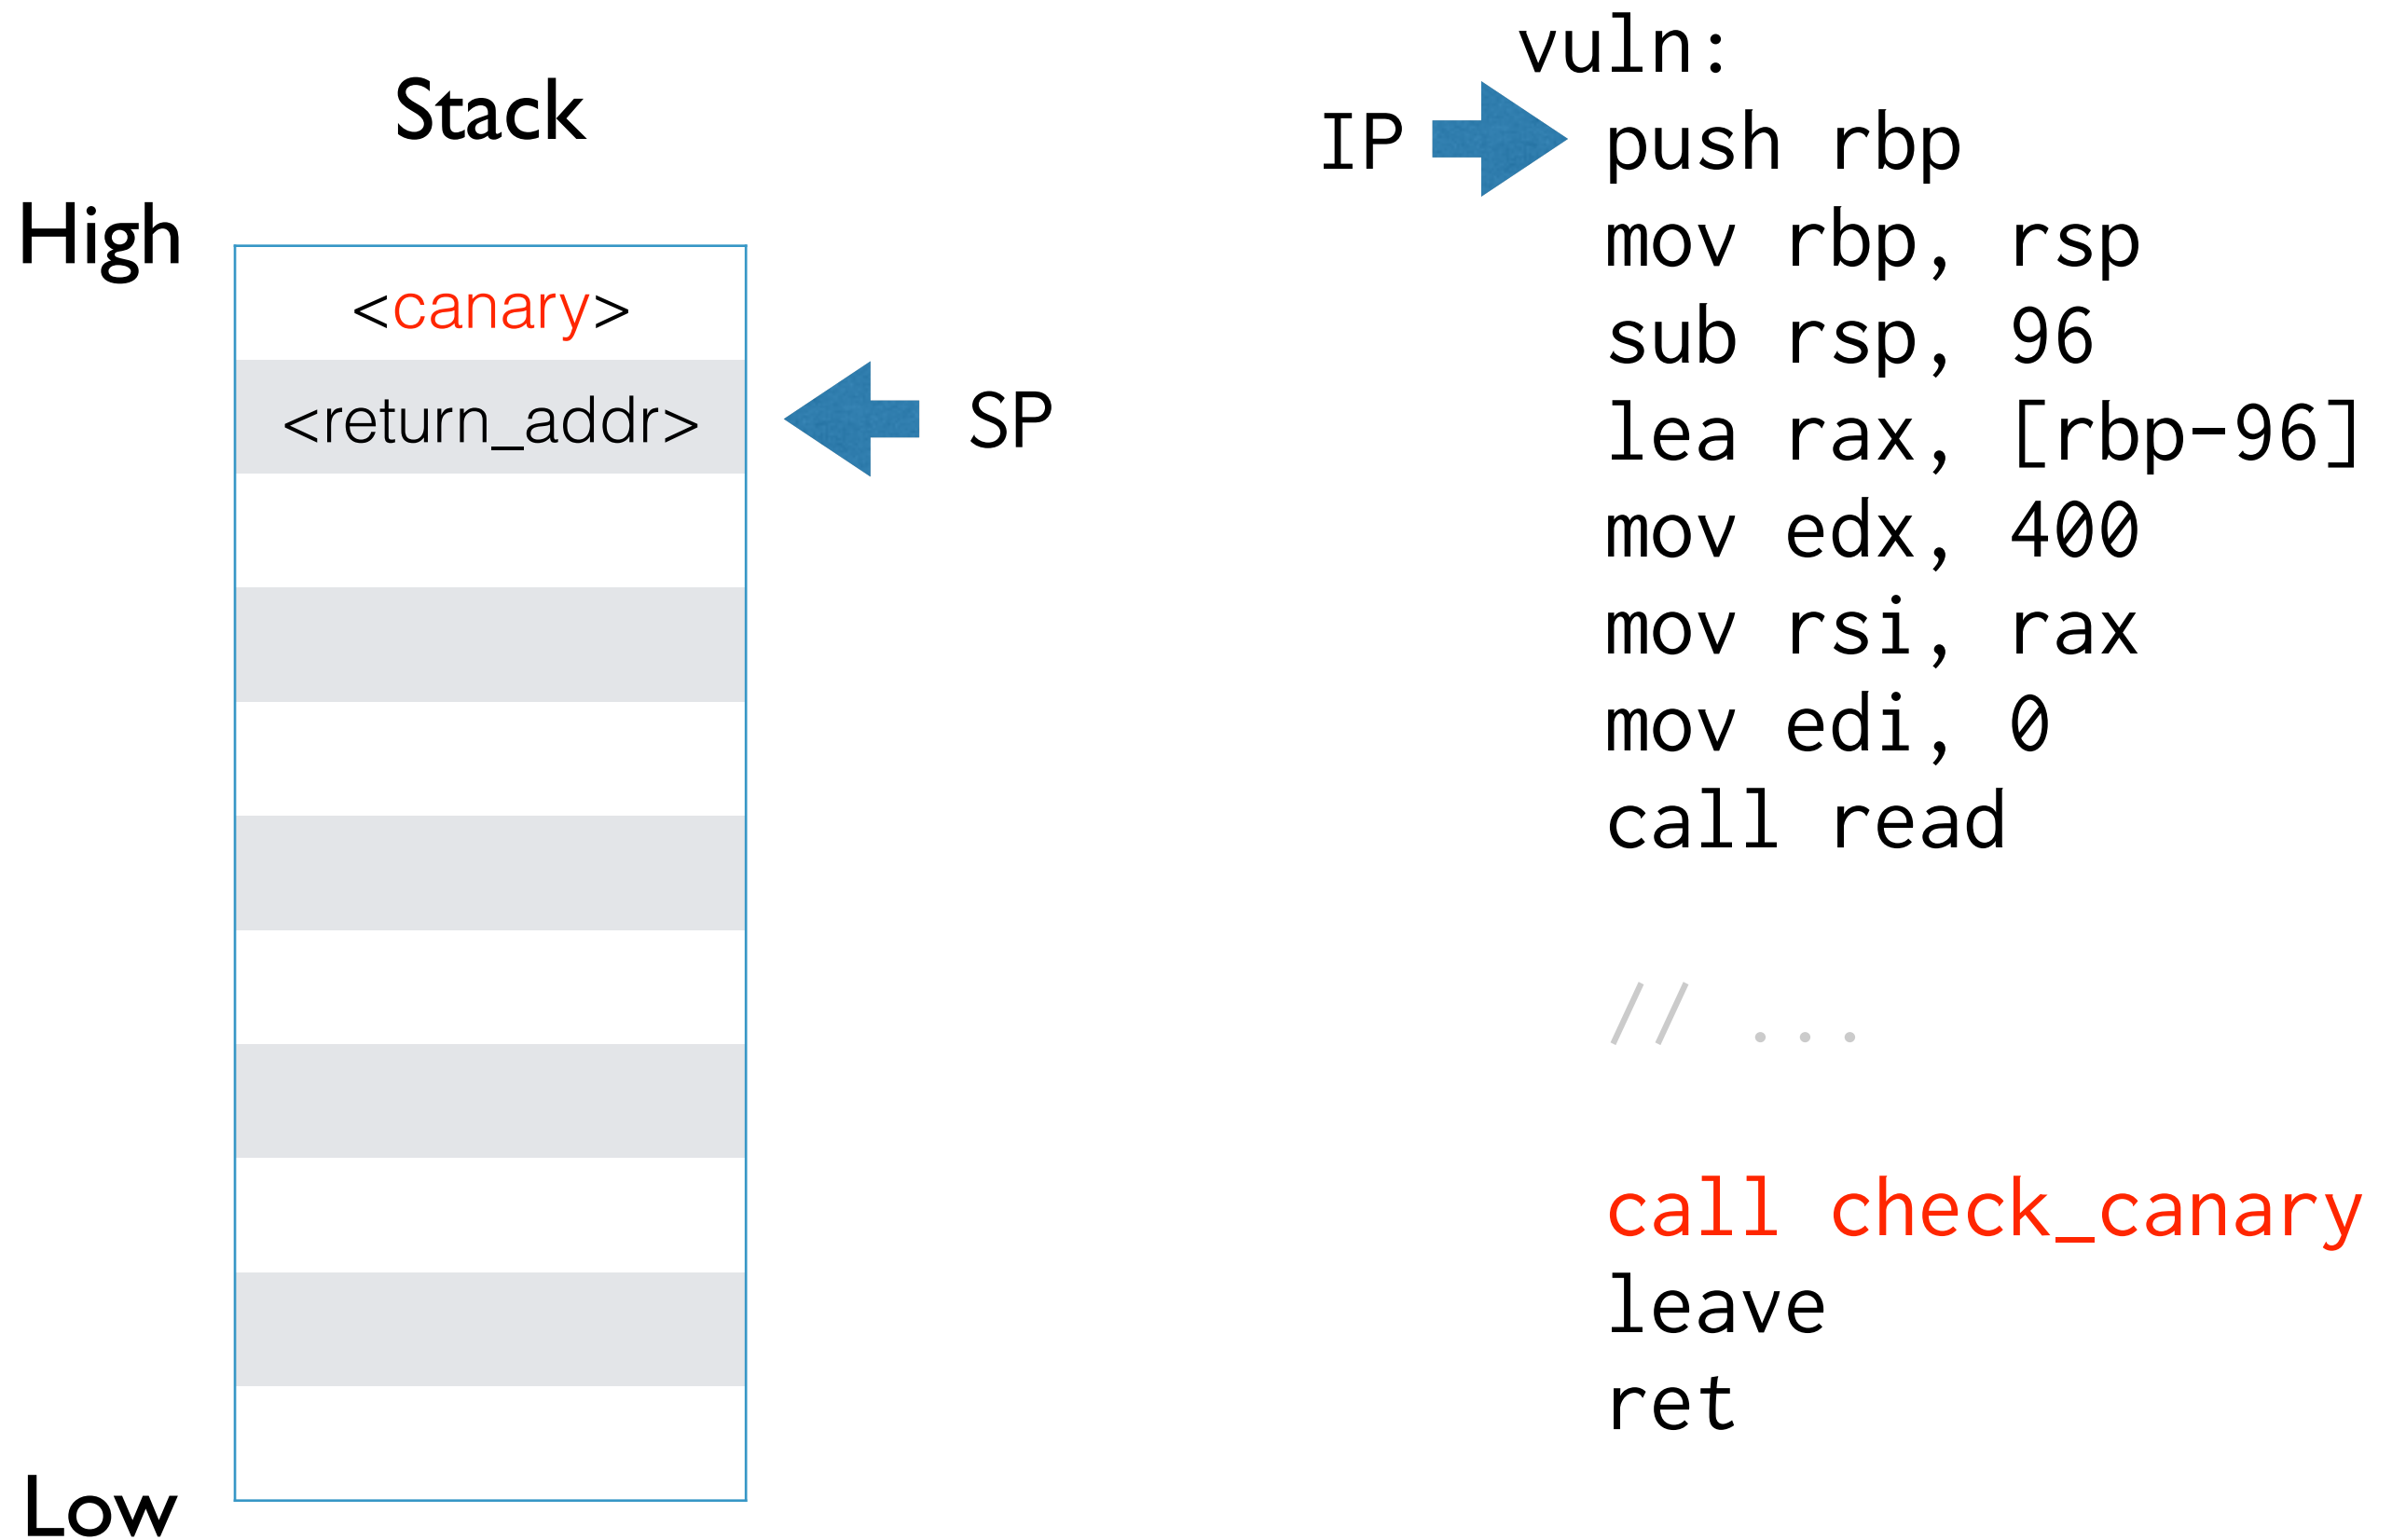
Shellcode, e.g.,
execve("/bin/sh")

as a string of bytes
\x31\xc0\x48 ...

Vulnerability (bug) +
Subvert Control Flow +
Payload (shellcode)

Defences

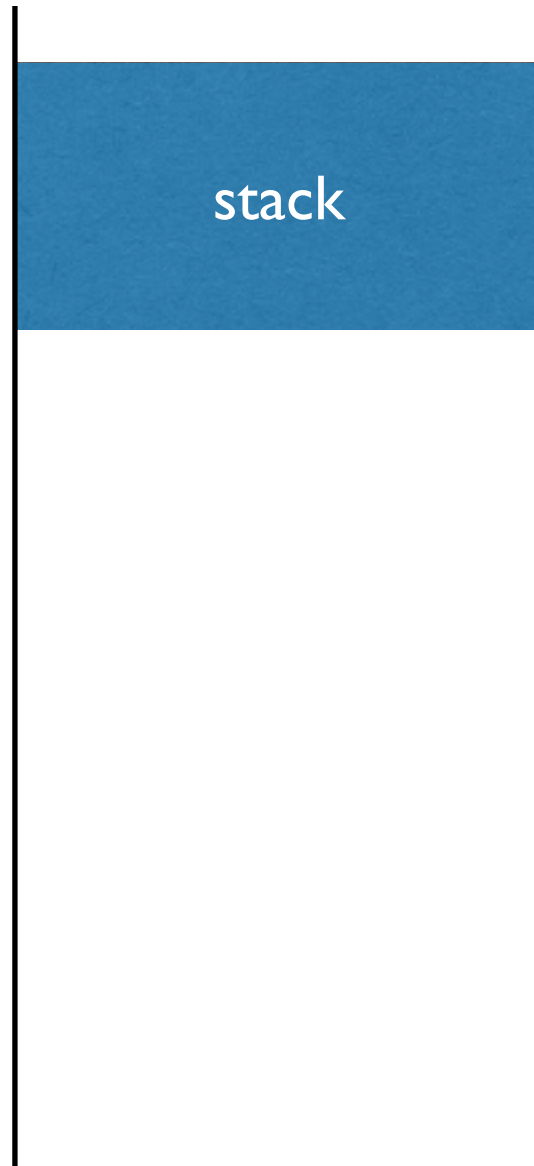
- Stack Canaries (remember `-fno-stack-protector`)
- Address Space Level Randomisation (ASLR)
- *Sanitise User Input*
- Execute Never (XN/W^X/DEP)



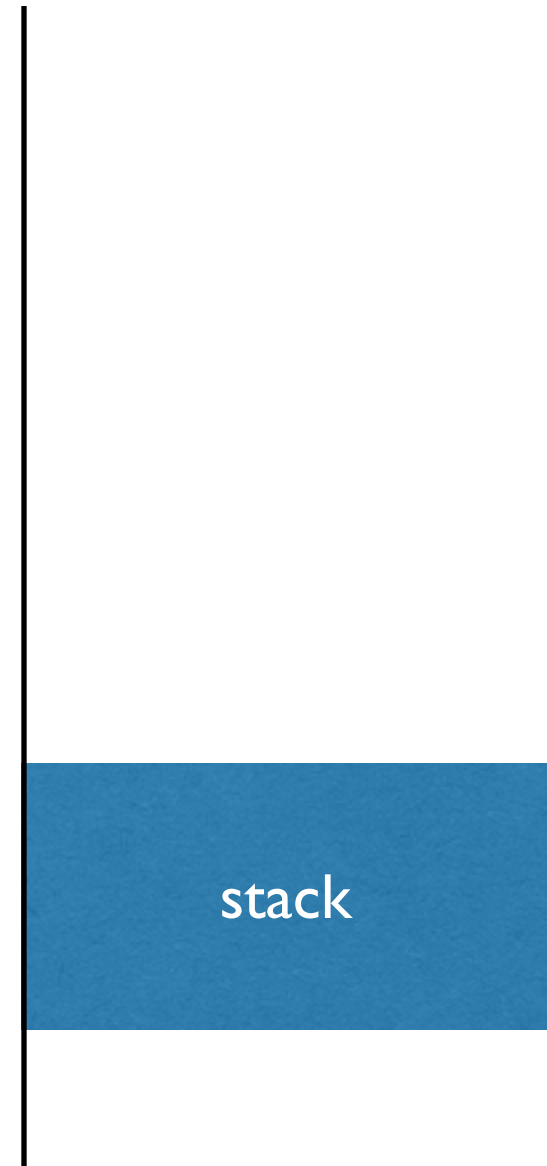
Run 0



Run 1



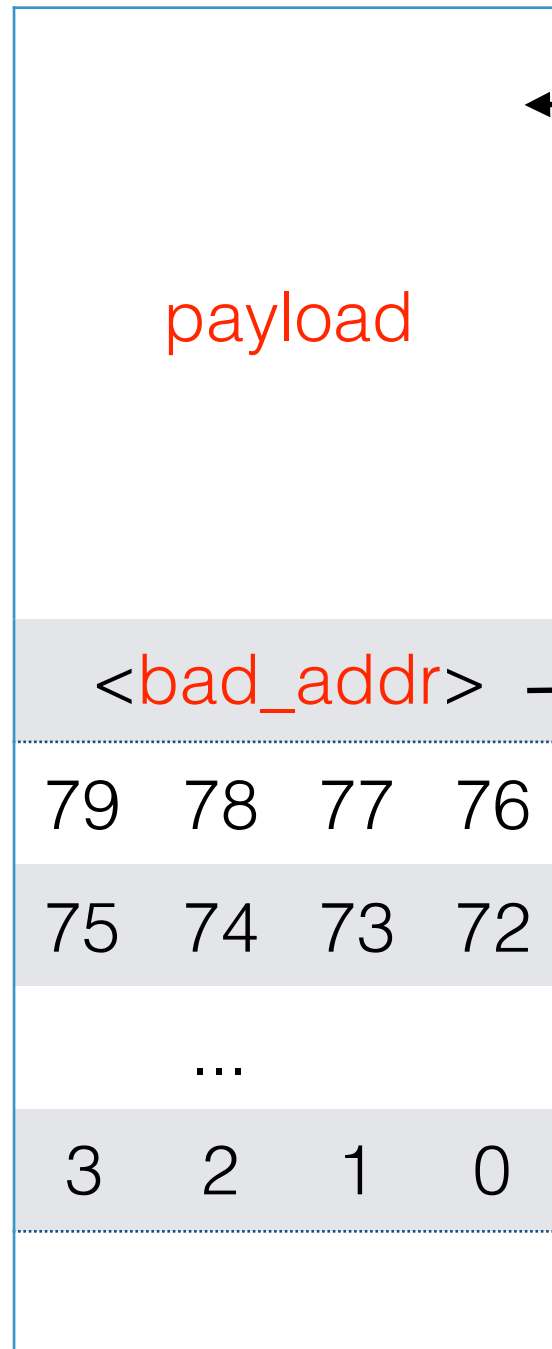
Run 2



run-to-run variation in layout of program

Stack

High



* mark stack (and other data areas)
as non-executable so this jump
causes a memory protection fault

Low

Stack Canaries

Vulnerability (bug) +
Subvert Control Flow +
Payload (shellcode)

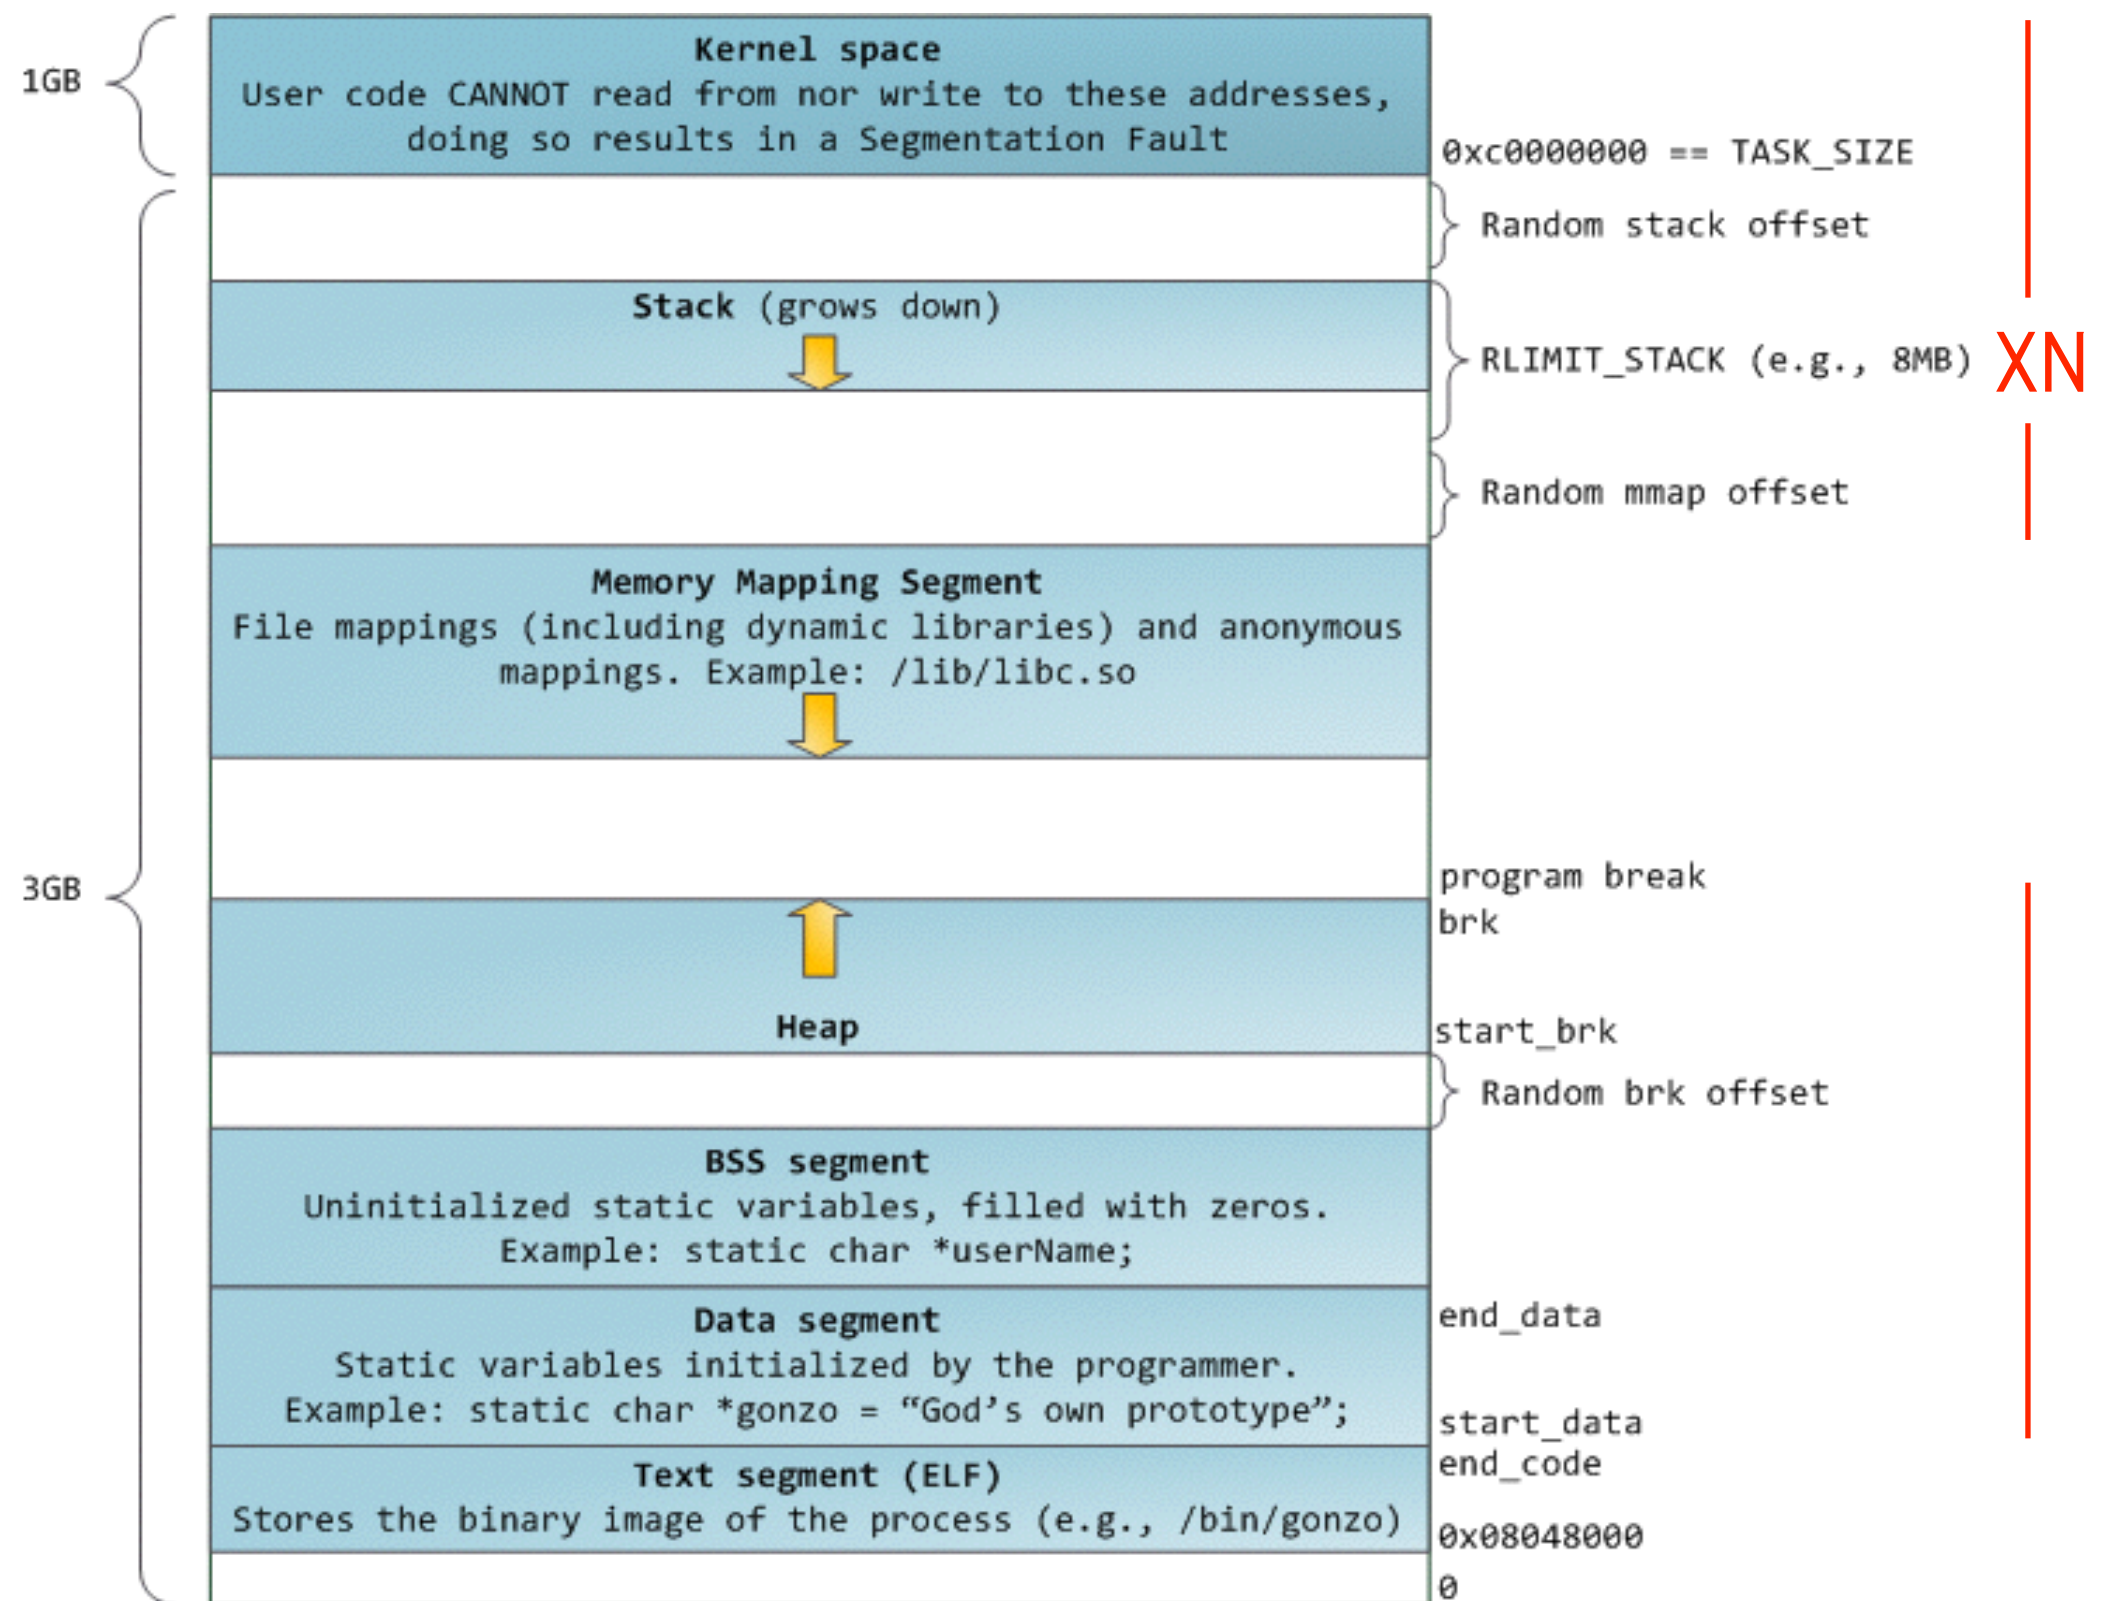
ASLR

Sanitise input, XN

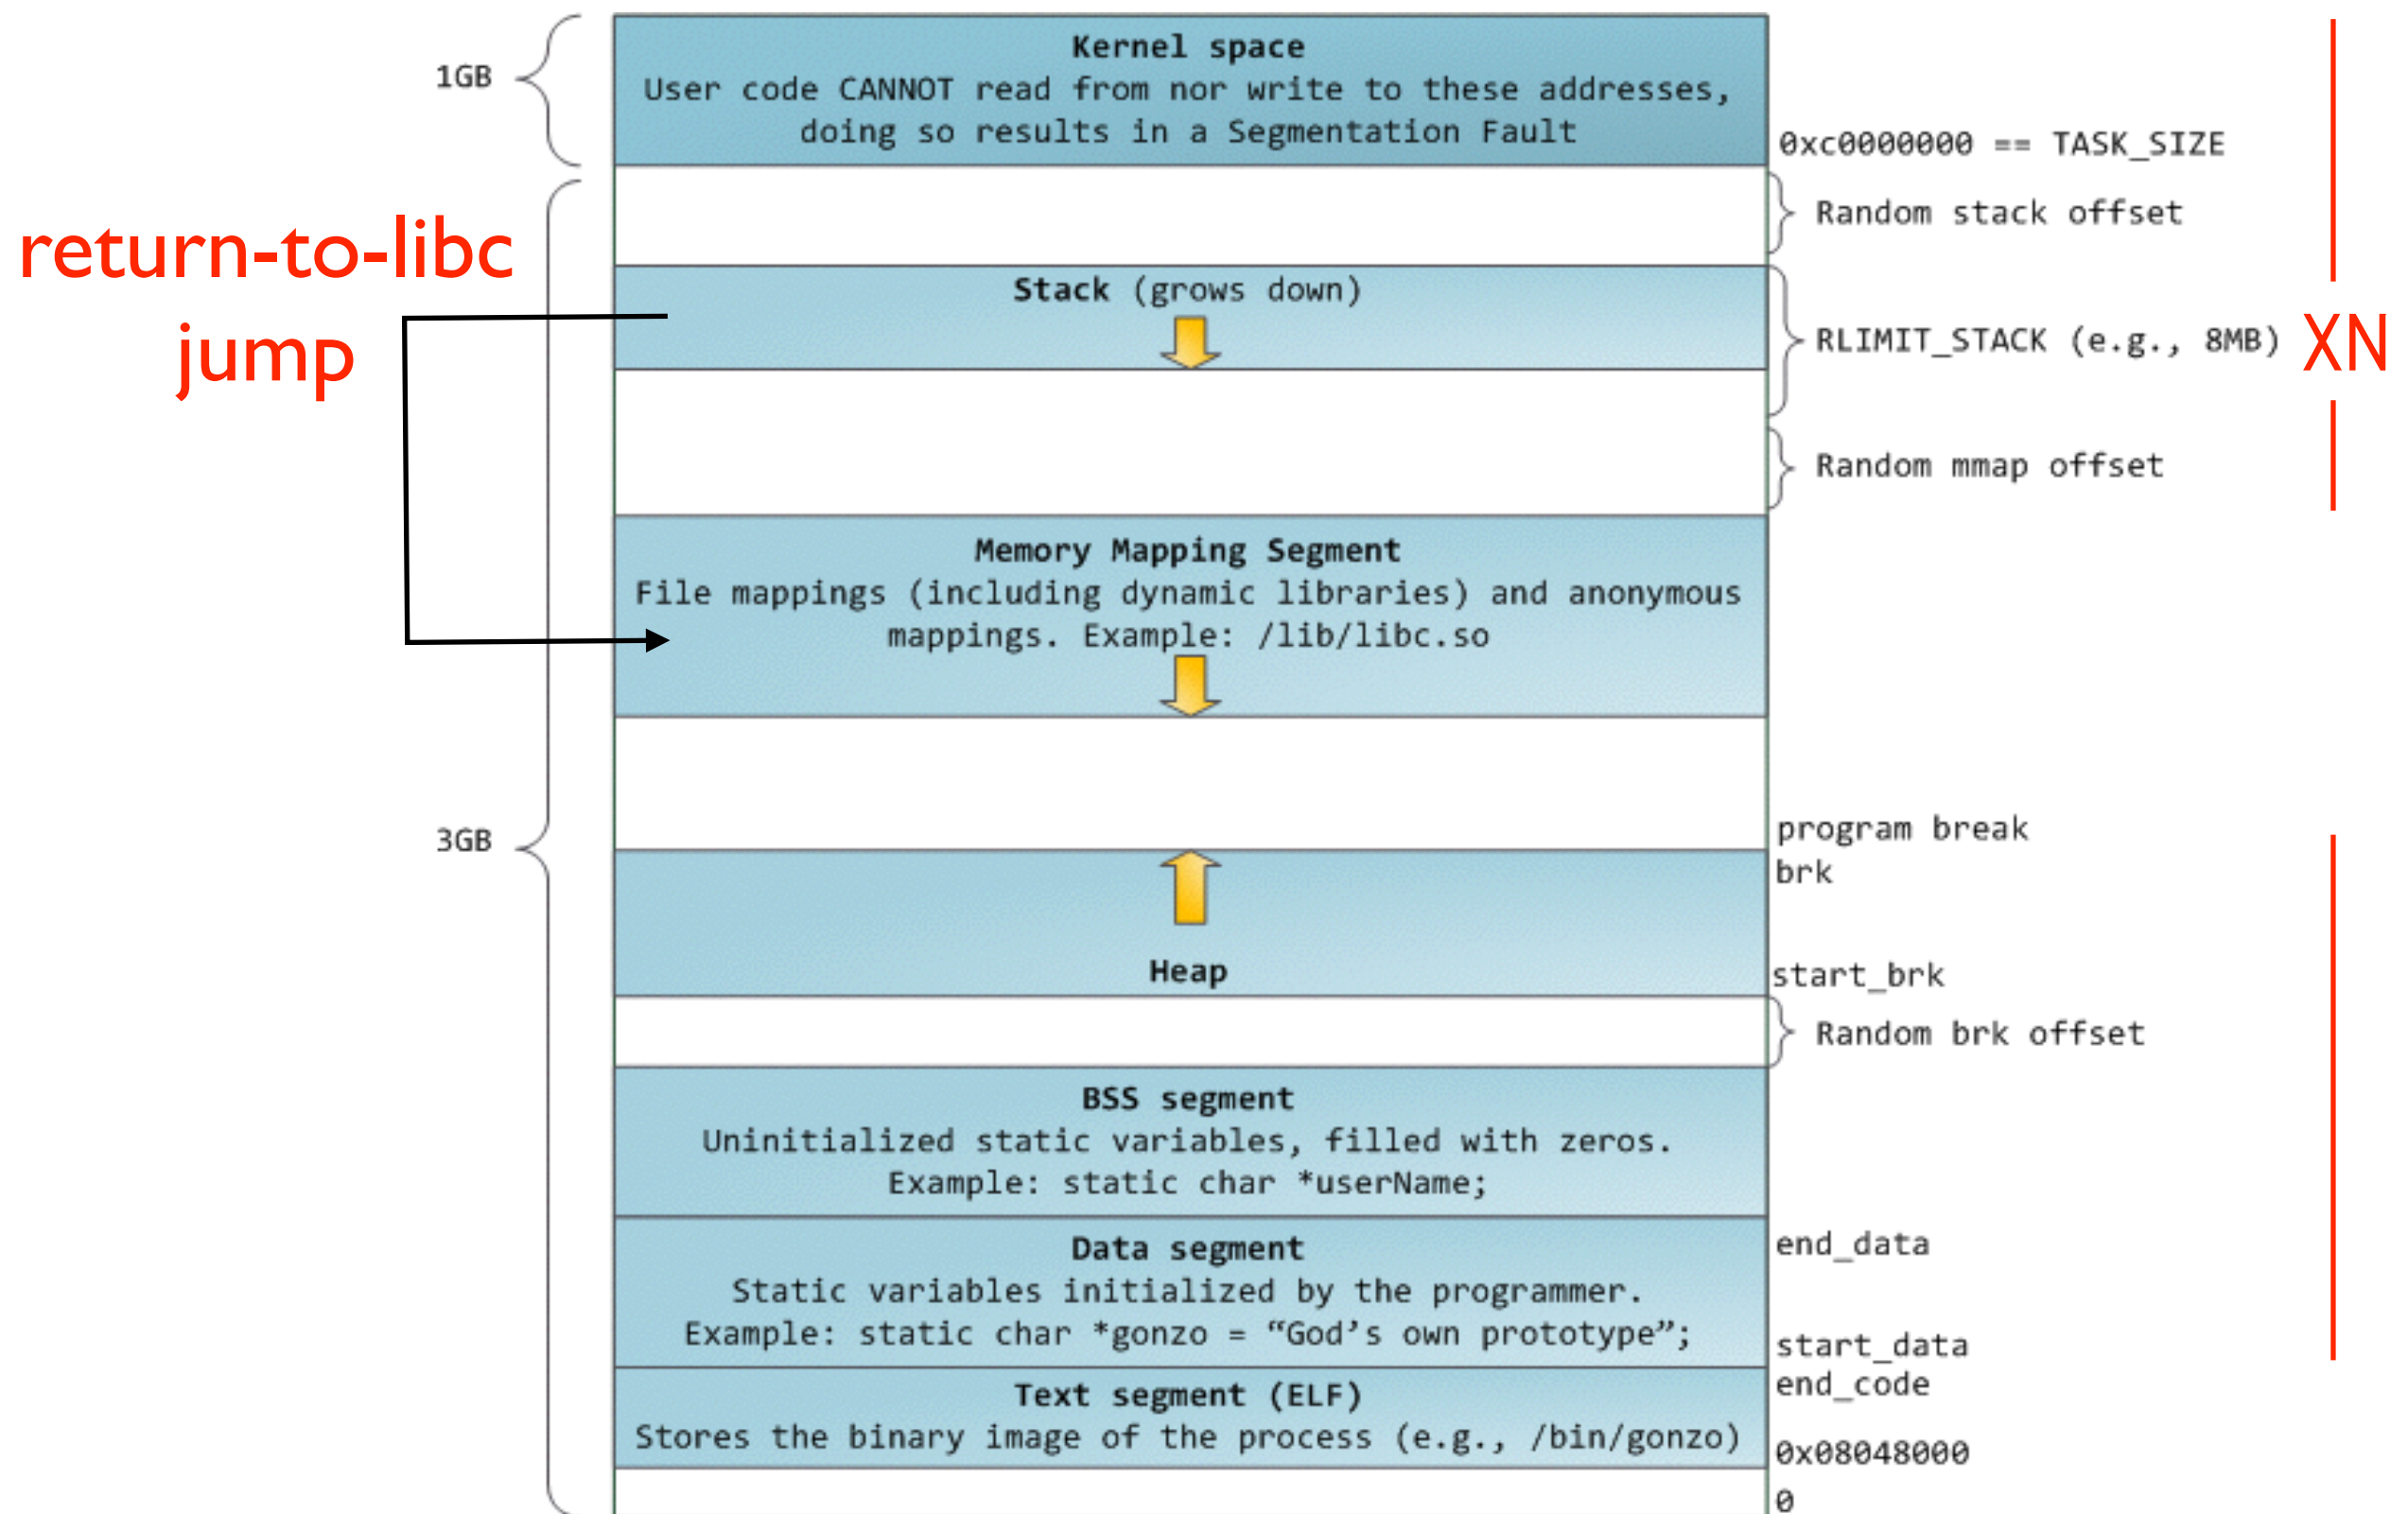
When?

- Programming (**test and verify**, input sanitisation)
- Compile-time / Dynamic Instrument (stack canaries)
- Runtime (ASLR, XN)

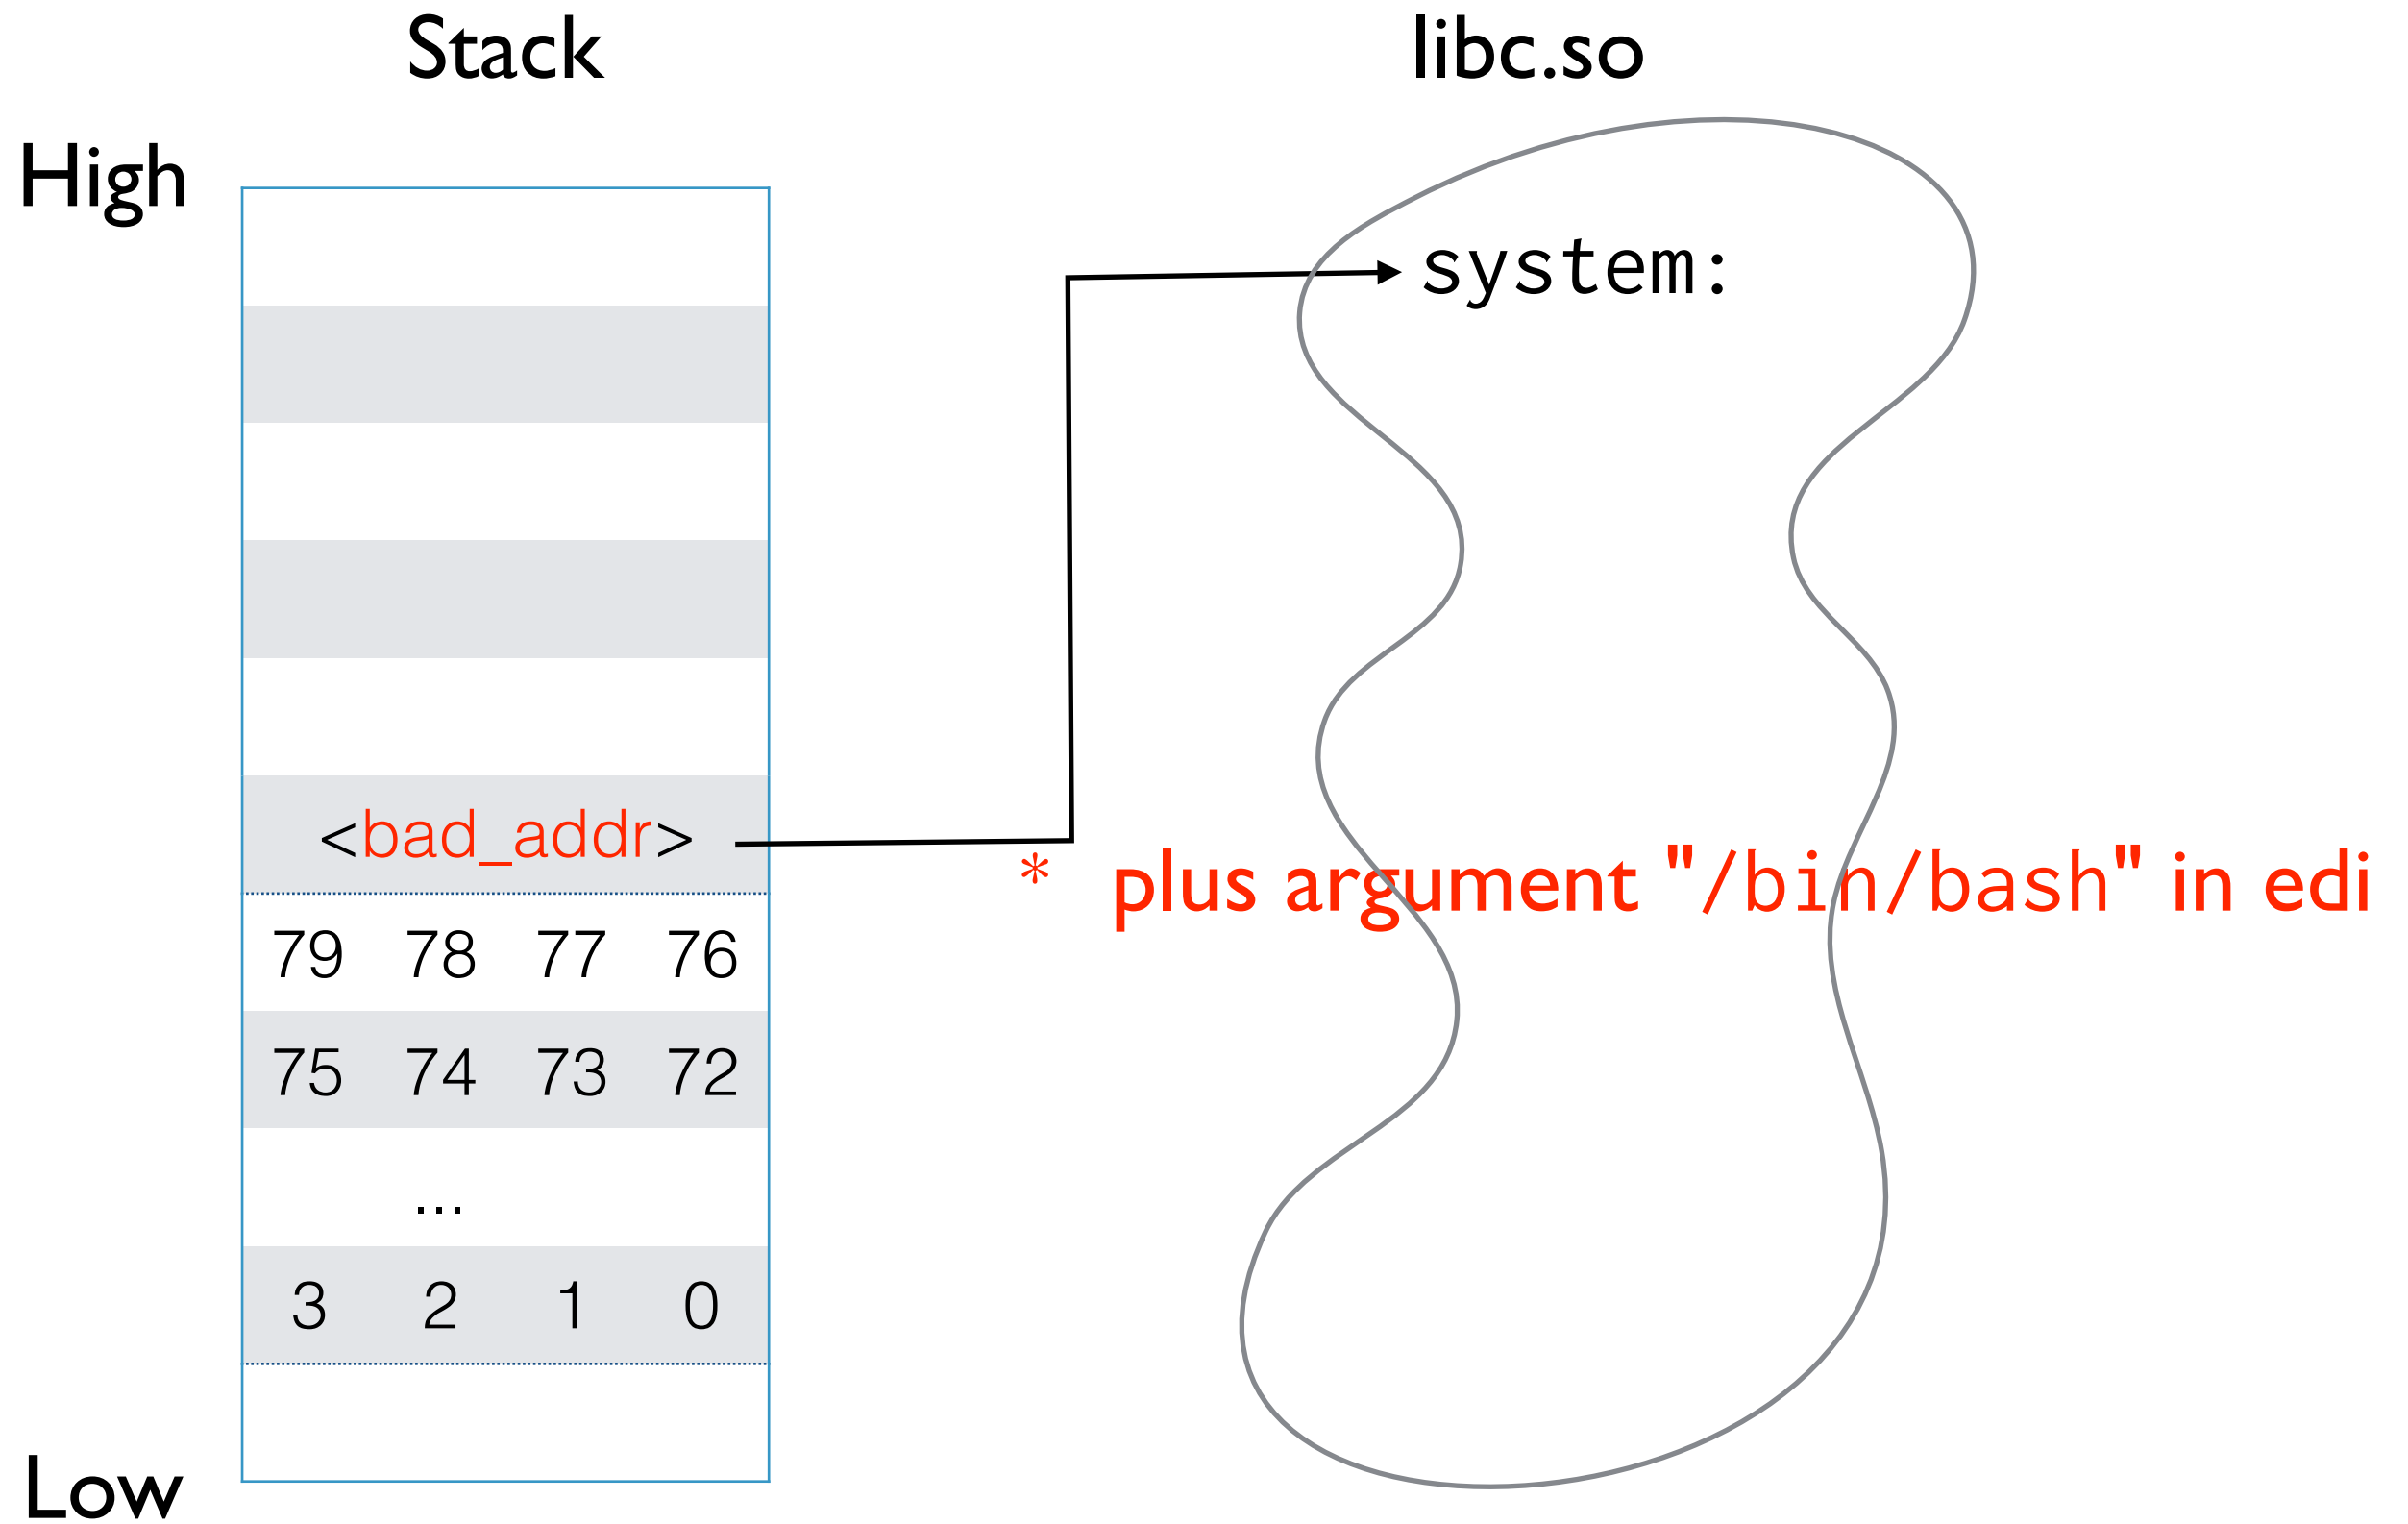
Return-to-Libc



<http://static.duarte.org/img/blogPosts/linuxFlexibleAddressSpaceLayout.png>



<http://static.duarte.org/img/blogPosts/linuxFlexibleAddressSpaceLayout.png>



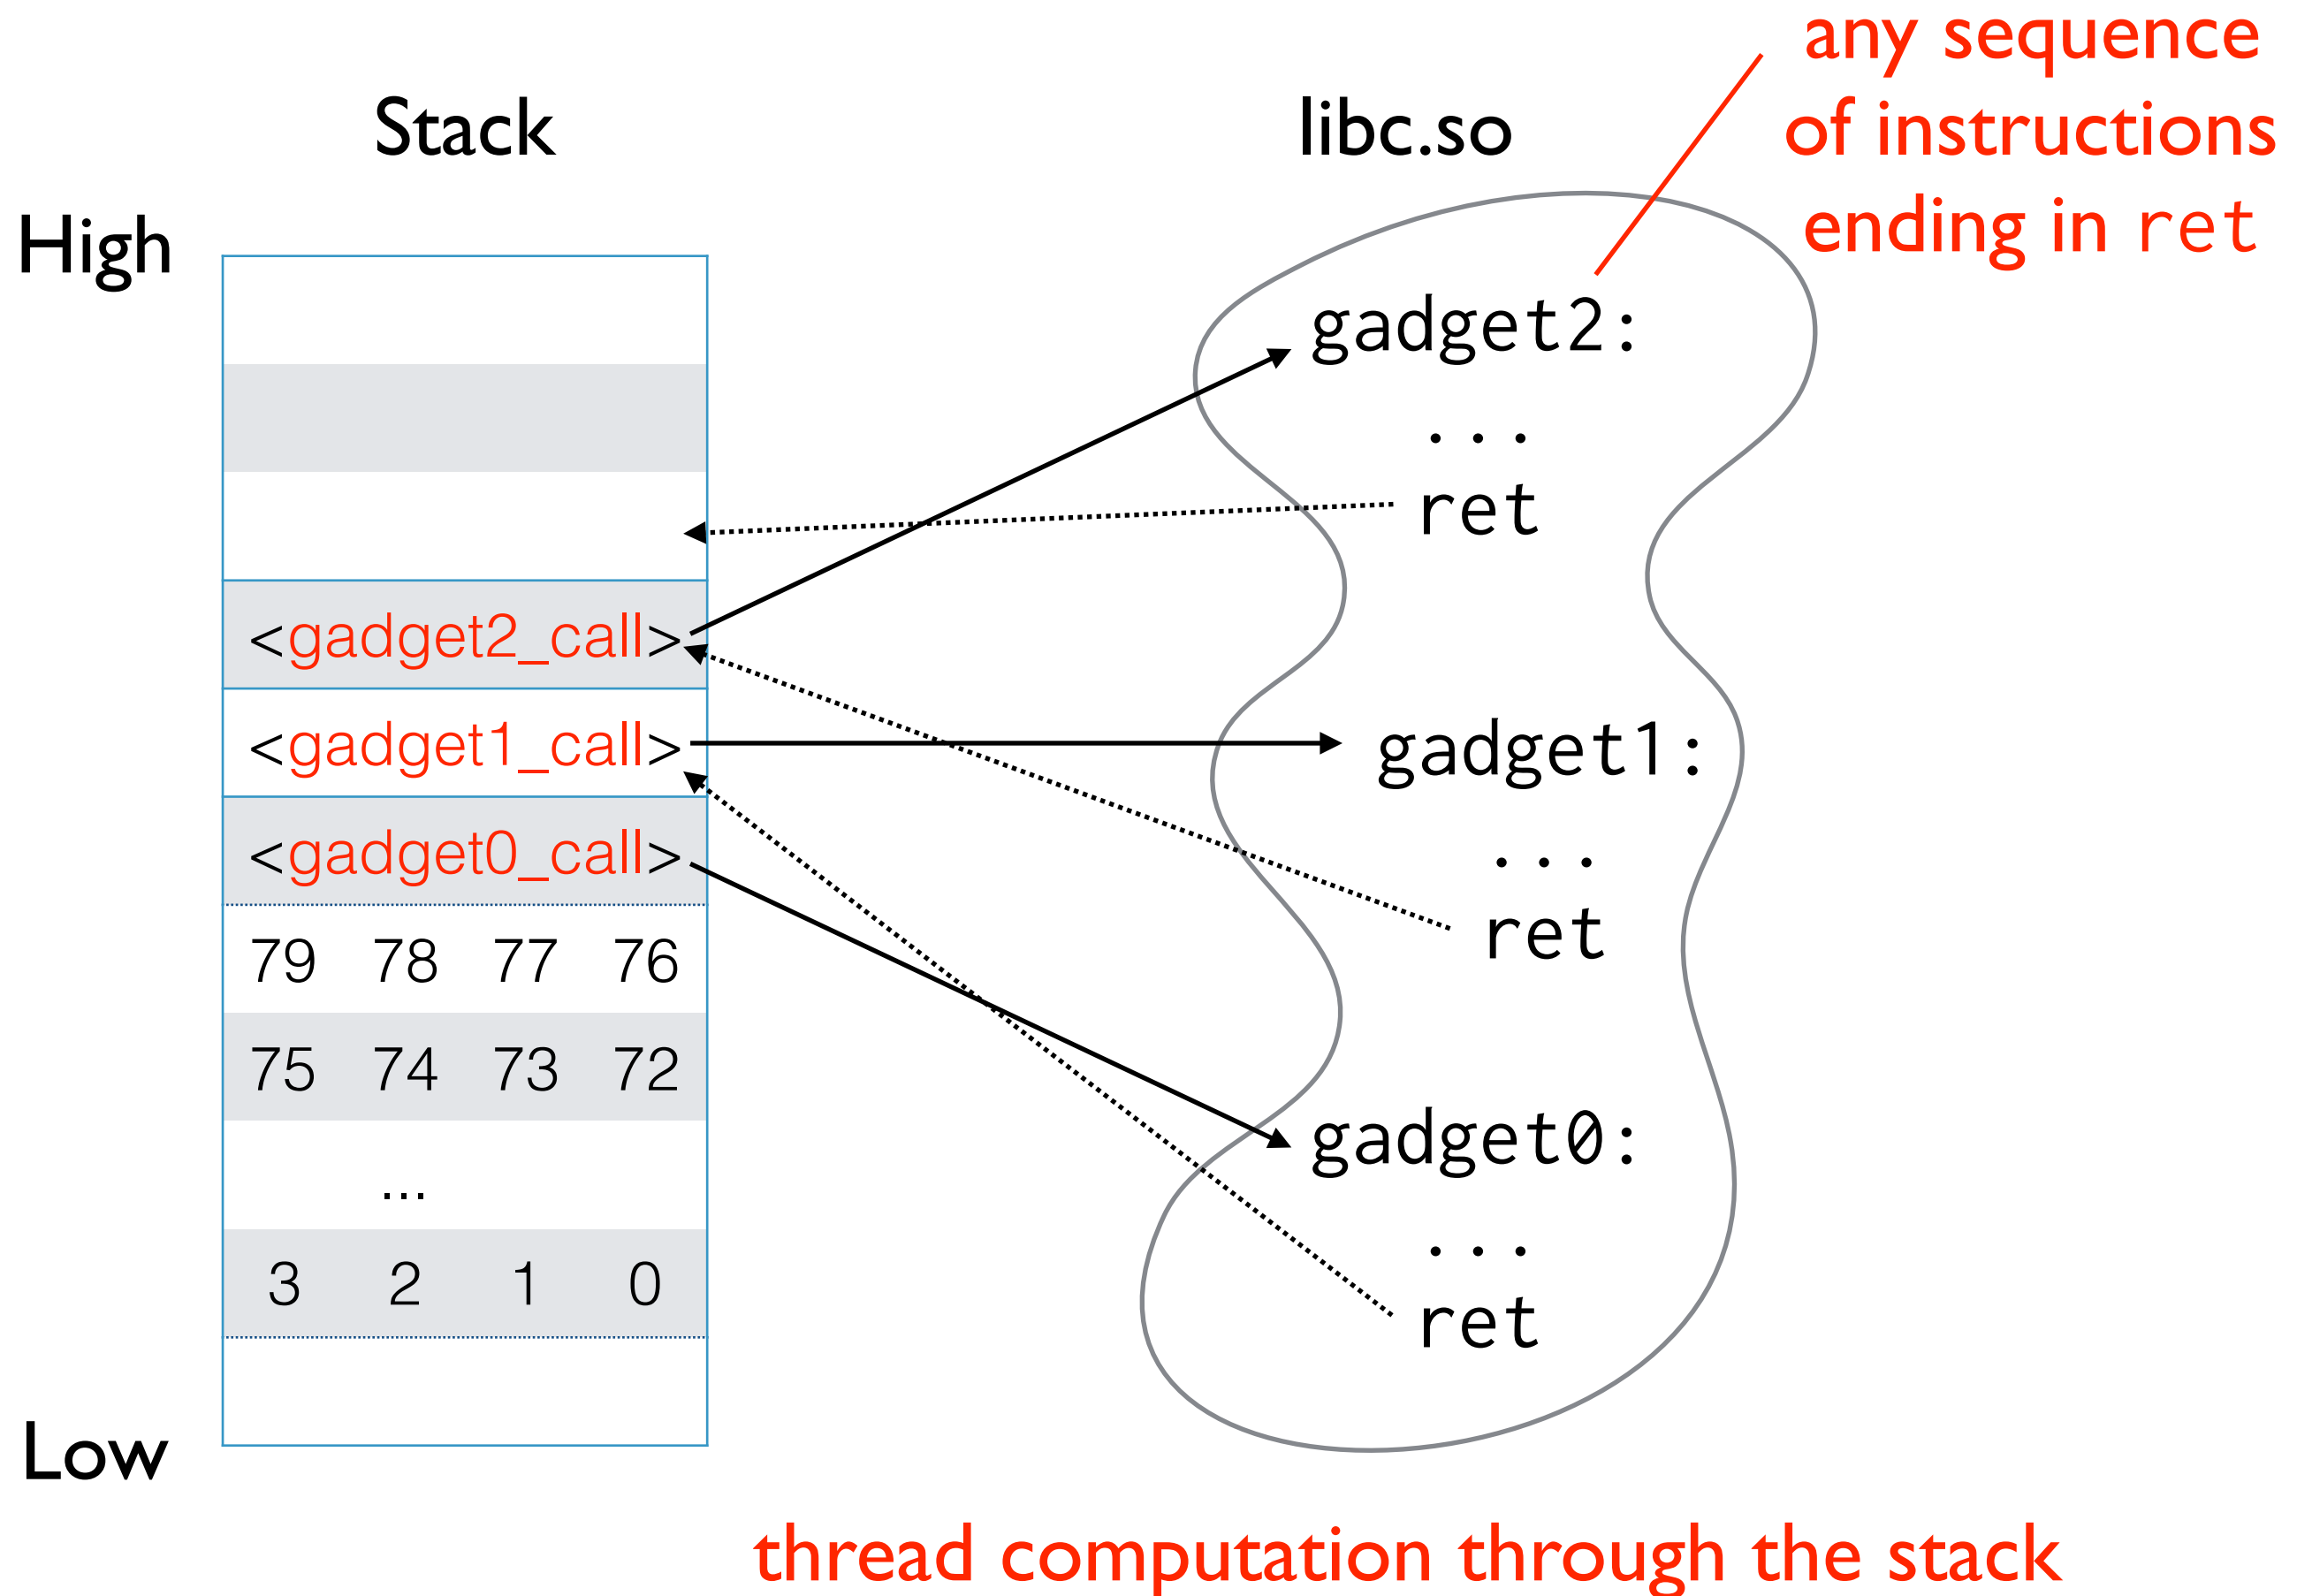
Assumed Limitations

- Straight-line calls into functions (no branching)
- Restricted to code in loaded libraries / text segment

Full version of an extended abstract published in Proceedings of ACM CCS 2007, ACM Press, 2007.

The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)

Hovav Shacham*
hovav@cs.ucsd.edu



Key Observation

```
f7 c7 07 00 00 00
0f 95 45 c3
```

```
test $0x00000007, %edi
setnzb -61(%ebp)
```

```
c7 07 00 00 00 0f
95
45
c3
```

```
movl $0x0f000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret
```

Our thesis: In any sufficiently large body of x86 executable code there will exist sufficiently many useful code sequences that an attacker who controls the stack will be able, by means of the return-into-libc techniques we introduce, to cause the exploited program to undertake arbitrary computation.

Gadgets

- NOP
- Load Immediate and Load / Store Memory
- Arithmetic: Add, Sub, Xor, Shifts, Rotates
- Control Flow: Jumps, System Calls, Function Calls
- Found in unintended sequences of libc (thwart XN)
- Allowing arbitrary computation (thwart limitations of ret2libc)

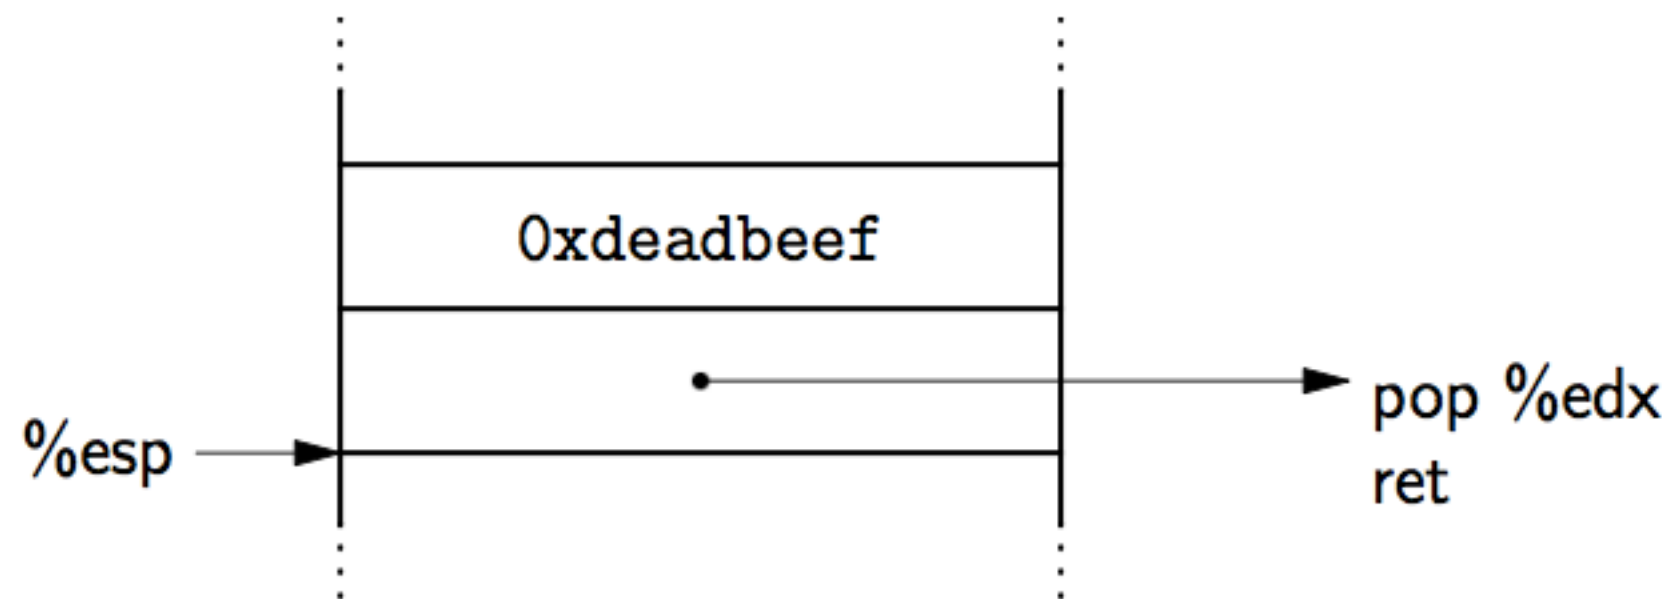


Figure 2: Load the constant 0xdeadbeef into %edx.

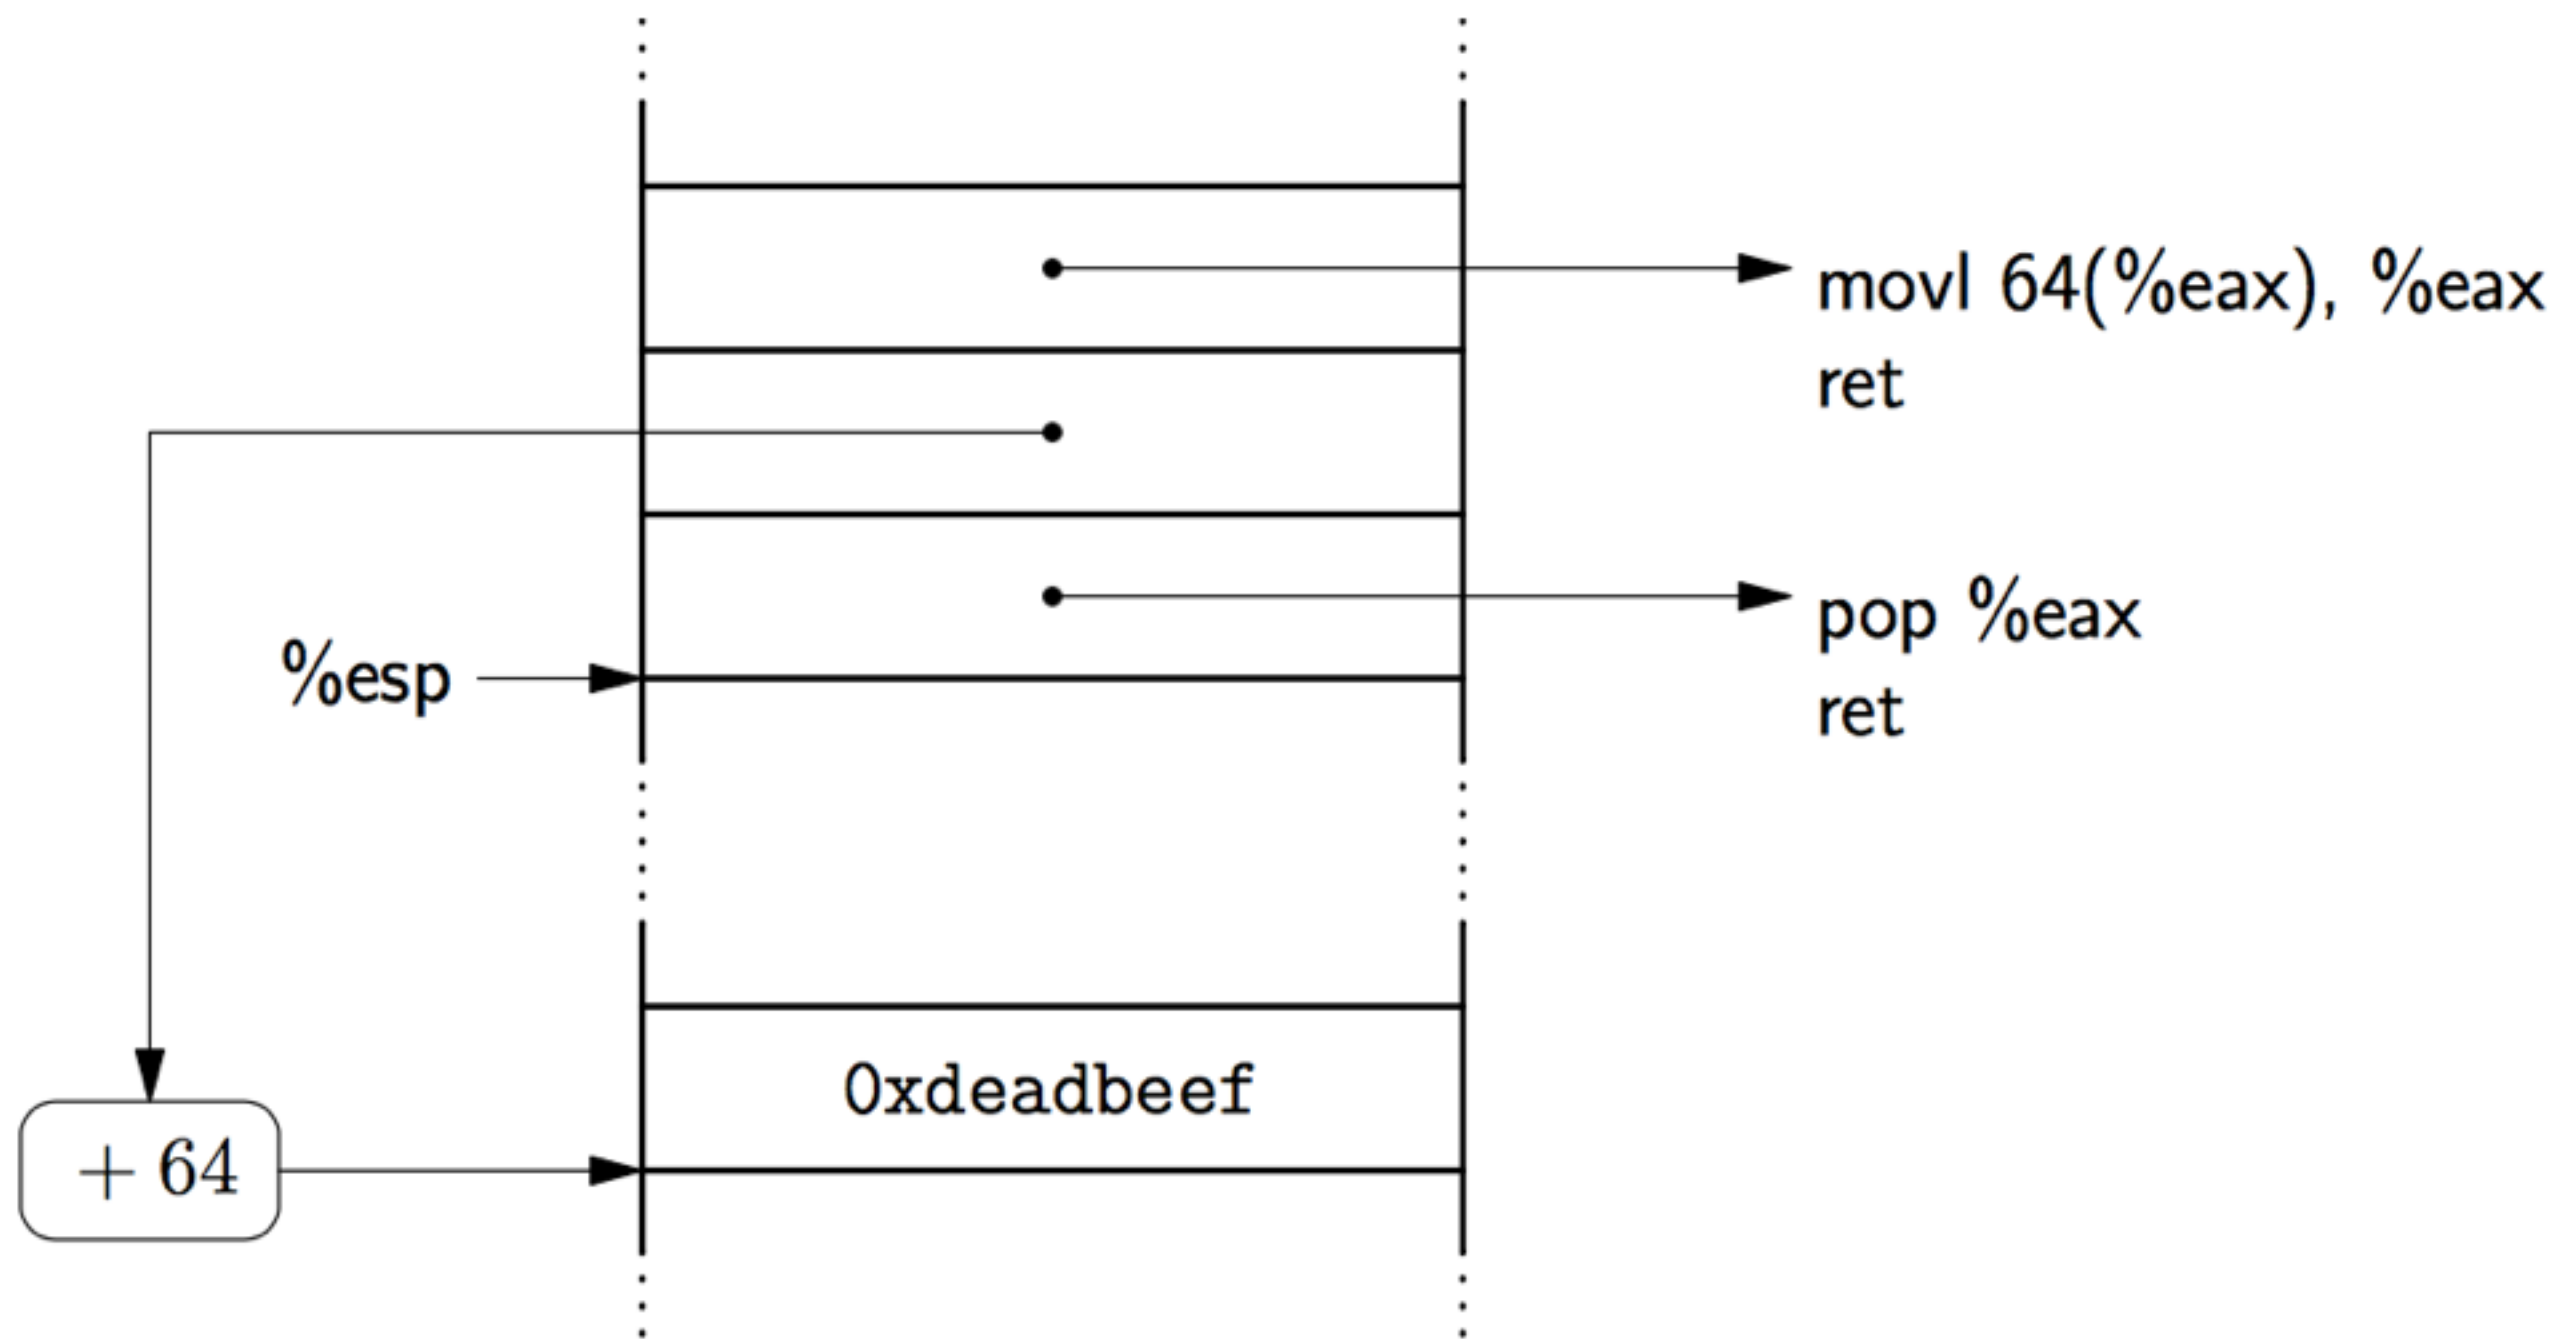


Figure 3: Load a word in memory into `%eax`.

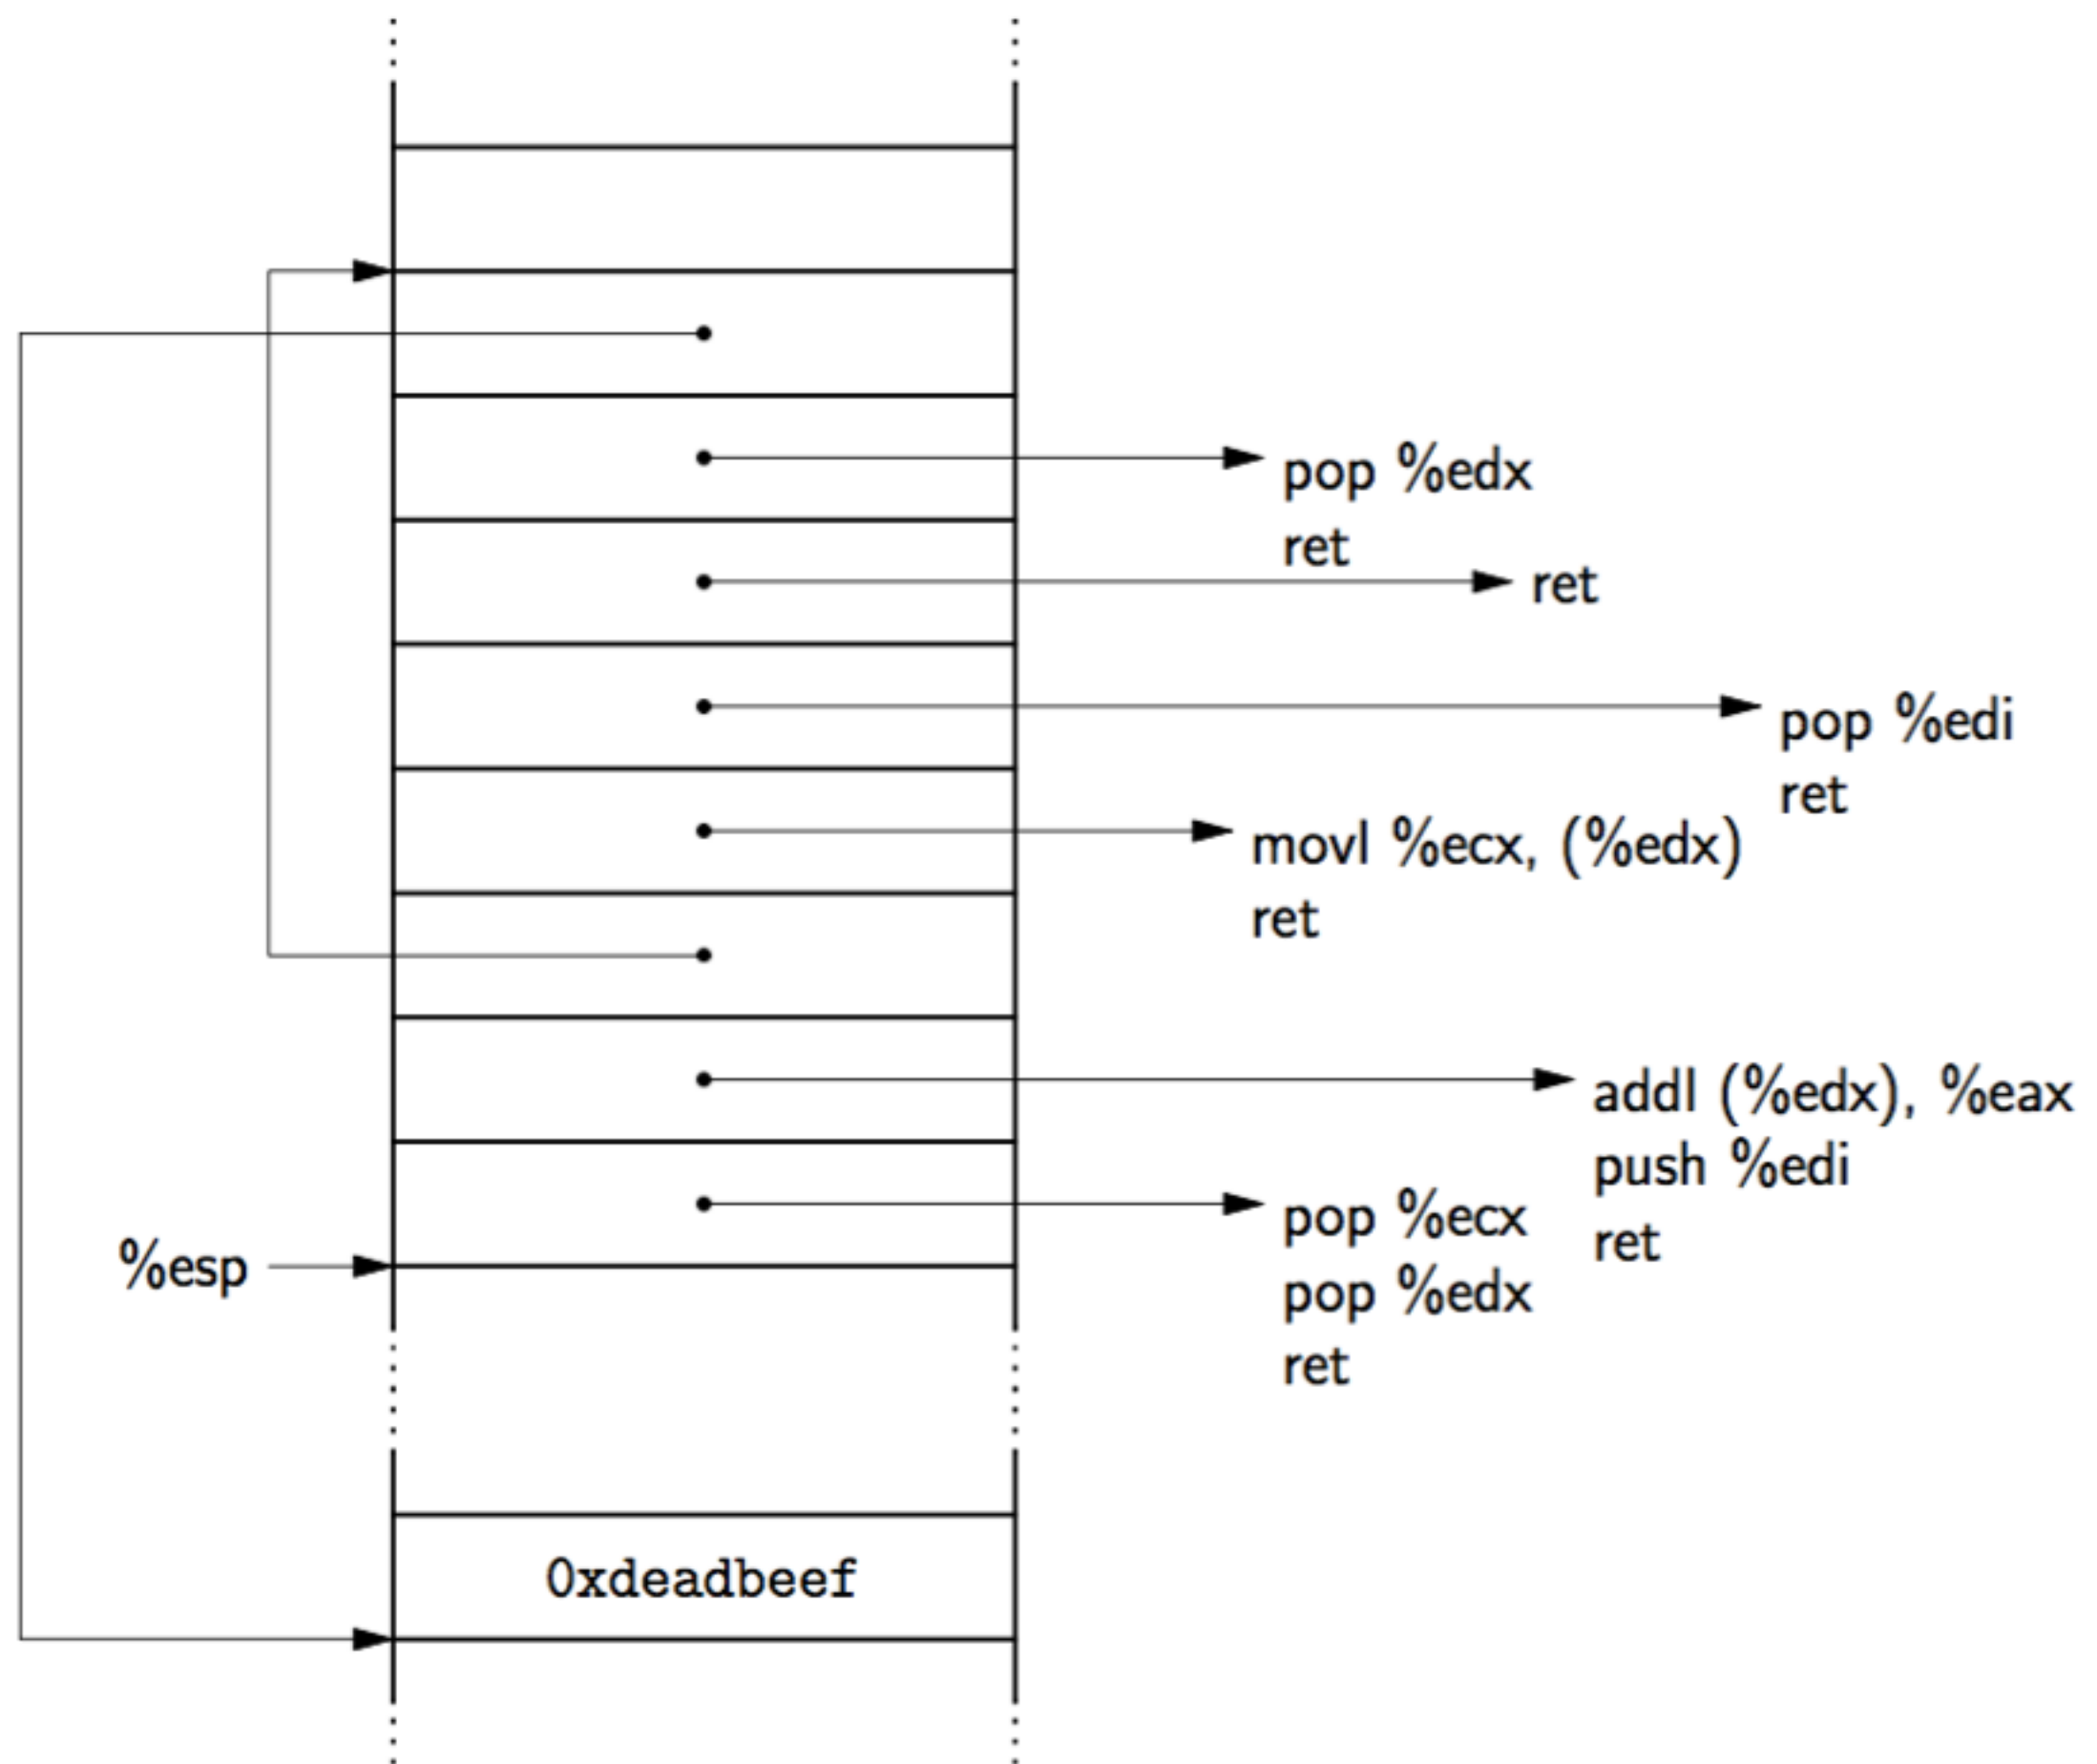


Figure 7: Repeatable add into `%eax`.

More precisely

- Search a given libc version for "interesting" sequences
- Construct gadgets for necessary operations
- Program a shellcode (e.g., `execv('/bin/bash')`) using this gadget programming language

Summary

- Interesting interplay between software testing/verification and security
- Arms race between attackers and defenders
- New architectural extensions to combat return-oriented programming